

CLASS III HYPERCUBOID CHARACTER SUMS AND A DISCRIMINANT-4 SINGULAR K3 SURFACE

ELIAS DE JESÚS

ABSTRACT. We study the character sum $\Sigma_{\text{III}}(p) = \sum_{x \in \mathbb{F}_p^3} \chi(x_0 x_1 x_2 (x_0 + x_2)(x_1 + x_2)(x_0 + x_1 + x_2))$, the third of three residual classes arising from a clean finite-field hypercuboid residue system, and previously left open where the companion Classes I/II were resolved. An elementary projective reduction gives $\Sigma_{\text{III}}(p) = (p-1)T(p)$ for the two-variable sum $T(p) = \sum_{x,t \in \mathbb{F}_p} \chi(xt(t+1)(x+t)(x+t+1))$, and an explicit coordinate involution proves $\Sigma_{\text{III}}(p) = 0$ for $p \equiv 3 \pmod{4}$ by an elementary, self-contained argument requiring no arithmetic geometry. For $p \equiv 1 \pmod{4}$, $T(p)$ is realized as the trace of Frobenius on the transcendental lattice of a singular K3 surface X_{III} , constructed as the elliptic fibration attached to $T(p)$'s defining affine surface. This surface has three fibers of Kodaira type I_2^* and full rational 2-torsion $(\mathbb{Z}/2)^2$ over $\mathbb{Q}(T)$, the field of the fibration's base. We prove that every irreducible component of every reducible fiber is individually defined over $\mathbb{Q}$ by a new argument – not requiring an explicit resolution of any singular fiber – combining the rationality of the zero section and the three nonzero 2-torsion sections with a rigidity property of the fiber's incidence graph: once the zero section and the three torsion sections are shown to occupy four pairwise distinct multiplicity-one components, the graph automorphism group compatible with that marking is trivial, forcing the remaining components to be Galois-fixed as well. This yields $\text{Tr}(F_p \mid \text{NS}(X_{\text{III}})) = 20p$ unconditionally, and $\text{disc NS}(X_{\text{III}}) = 4$, $T(X_{\text{III}}) \cong \text{diag}(2, 2)$, identifying the associated CM field as $\mathbb{Q}(i)$. Combining this with Livné's modularity theorem and a finite twist-elimination argument (exhaustive over a candidate set of four characters, forced by ramification, before any numerical value is consulted), we prove $\text{Tr}(F_p \mid T(X_{\text{III}})) = a_p(\mathbf{16.3.c.a})$ for every odd p , where $\mathbf{16.3.c.a} = \eta^6(4z)$ is Ahlgren–Ono–Penniston's own modular form, not claimed as new here; combined with an independently re-derived point-count ledger, this yields the exact evaluation $\Sigma_{\text{III}}(p) = (p-1)a_p(\mathbf{16.3.c.a})$ for $p \equiv 1 \pmod{4}$ and $\Sigma_{\text{III}}(p) = 0$ for $p \equiv 3 \pmod{4}$, completing the finite-field hypercuboid classification left open by the companion paper. We also record the surface's provenance: $T(p) = A(-1, p)$ for Ahlgren–Ono–Penniston's two-variable sum, and X_{III} is isomorphic over $\mathbb{Q}$ to the $\lambda = -1$ member of their family, whose transcendental lattice $\text{diag}(2, 2)$, discriminant 4 and CM field $\mathbb{Q}(i)$ are already recorded in the literature. The hypercuboid route therefore reaches a known singular-K3 arithmetic sector independently rather than exhibiting a new one; what is contributed here is the combinatorial origin, the finite-field identities, and the rationality argument. We note that $\lambda = -1$ is precisely the parameter at which Ahlgren–Ono–Penniston's own closed-form evaluation degenerates, so the evaluation obtained here is not a specialization of theirs.

1. INTRODUCTION

This paper is a sequel to *Finite-Field Hypercuboid Character Sums and a Discriminant-8 Singular K3 Surface* [1], which studied two of the three residual character-sum classes, $\Sigma_I(p)$ and $\Sigma_{\text{II}}(p)$, arising from the zero-diagonal complement collapse structure of a clean finite-field hypercuboid residue system on four coordinates. That paper reduced $\Sigma_I(p)$ to a two-variable sum realized, via an explicit singular K3 surface, in terms of the weight-3, level-8 newform $f = \mathbf{8.3.d.a}$, with CM by $\mathbb{Q}(\sqrt{-2})$ and transcendental lattice $T(X) \cong \text{diag}(2, 4)$. The present paper resolves the third class, $\Sigma_{\text{III}}(p)$, whose elementary vanishing for $p \equiv 3 \pmod{4}$ was already established in [1] but whose value for $p \equiv 1 \pmod{4}$ was left as the genuine open problem of that paper.

We show that $\Sigma_{\text{III}}(p)$, for $p \equiv 1 \pmod{4}$, is governed by a *different* singular K3 surface X_{III} , with discriminant 4 and CM by $\mathbb{Q}(i)$ – structurally distinct from the discriminant-8, $\mathbb{Q}(\sqrt{-2})$ surface of [1]. We emphasize at the outset that this is an *observation about the present hypercuboid*

system, not a general theorem: we do not claim, and this paper does not attempt to prove, that the hypercuboid classification’s constraint classes always select distinct CM fields, or that any such correlation persists beyond the two branches exhibited here. Section 9 states this explicitly alongside a comparison table.

The main methodological distinction from [1]. The two papers’ hardest technical steps differ in kind. In [1], the singular K3 surface has Mordell–Weil rank 1 over a quadratic extension, and the Picard-lattice argument proceeds via an explicit extra Mordell–Weil section together with an explicit twist isomorphism between two elliptic curves. Here, X_{III} already achieves the maximal geometric Picard number 20 from its trivial lattice alone (Mordell–Weil rank 0), and the open step is instead to show that the Néron–Severi lattice – which a priori need only be defined over $\overline{\mathbb{Q}}$ – is in fact defined over $\mathbb{Q}$. We resolve this not by resolving each of X_{III} ’s three singular fibers explicitly, but by an argument combining full rational 2-torsion with a rigidity property of the fiber’s incidence graph (Section 5). This is a different technique from the one used in [1], and is the paper’s central technical contribution. Its graph-rigidity step is a standard corollary of published facts about the Galois action on I_n^* fibers (Remark 5.9); what is contributed is the assembly – deducing NS-rationality from rational 2-torsion alone, with no fiber resolved.

Relation to Ahlgren–Ono–Penniston. The governing modular form identified in this paper’s later sections, LMFDB 16.3.c.a, is *not new*: it is exactly Ahlgren, Ono, and Penniston’s own $A(z) = \eta^6(4z)$ [2], the modular form attached to their singular K3 surface at their parameter $\lambda = 8$.

More than that, and more importantly for provenance, **the Class III surface is itself a member of Ahlgren–Ono–Penniston’s family.** Write $A(\lambda, p) = \sum_{x, y \in \mathbb{F}_p} \chi(xy(x+1)(y+1)(x+\lambda y))$ for the two-variable sum of [2], whose surfaces X_λ are the minimal resolutions of the double covers of $\mathbb{P}^2$ branched along the six lines $x, y, x+1, y+1, x+\lambda y$ and the line at infinity. Then

$$T(p) = A(-1, p) \quad \text{for every odd prime tested,}$$

and at the level of surfaces the branch sextic of X_{III} is carried *onto* that of the $\lambda = -1$ member, with no residual scalar, by the $\mathbb{Q}$ -linear substitution $(x, t, w) \mapsto (x, t-x, -t-w)$. Hence X_{III} is isomorphic over $\mathbb{Q}$ to X_{-1} . Both statements are checked in the script `class3_aop_identification.py`.

It is worth being exact about *which* member this is. [2] evaluate $A(\lambda, q)$ in closed form through a quantity of the shape $\chi(\lambda+1)(a(\lambda, q)^2 - q)$, which requires $\lambda+1$ to be invertible; van Geemen and Top likewise take $t \neq 0, -1$ in their isogeny theorem [3]. So $\lambda = -1$ is precisely the parameter at which the family’s standard evaluation degenerates – and it is the parameter at which the hypercuboid reduction lands. The same script confirms that the closed form reproduces the direct sum at $\lambda = 8$ and $\lambda = 2$ and is undefined at $\lambda = -1$.

The consequence must be stated plainly. The $\lambda = -1$ member is a classically studied surface: van Geemen and Top [3] record that its transcendental lattice has rank 2 and discriminant 4, hence is $\langle 2 \rangle \oplus \langle 2 \rangle$, attributing the lattice computation to Persson [4], and note that Shioda and Inose [12] exhibit the surface as the desingularised quotient of $E_i \times E_i$ with $E_i = \mathbb{C}/\mathbb{Z}[i]$ – that is, with complex multiplication by $\mathbb{Q}(i)$. **The arithmetic package $T(X) \cong \text{diag}(2, 2)$, disc NS = 4, CM by $\mathbb{Q}(i)$ is therefore already present in the literature and is not claimed here as a new discovery.** What this paper contributes is the route to it: the hypercuboid origin of the sum, the projective reduction $\Sigma_{\text{III}}(p) = (p-1)T(p)$, the self-contained involution proof of vanishing for $p \equiv 3 \pmod{4}$, the torsion-and-graph-rigidity proof that every fiber component is defined over $\mathbb{Q}$ without resolving any fiber, and the resulting exact evaluation of $\Sigma_{\text{III}}(p)$. The hypercuboid construction *independently reaches and realizes* a known singular-K3 arithmetic sector; it does not discover one.

Two cautions about the $\lambda = 8$ comparison, which earlier drafts of this paper rested on. The surfaces do have different branch configurations: $\lambda = 8$ lies in the generic incidence stratum of

the family (three triple points and six double points), whereas $\lambda = -1$ lies in the stratum with four triple points and three double points. But the observation that $T(p)$ and $A(8, p)$ differ is *not* by itself evidence that the underlying surfaces are non-isomorphic: the difference is the entirely structured

$$A(8, p) - T(p) = \left(\frac{-1}{p}\right)p,$$

a term of exactly the size produced by a change of affine normalisation rather than by a change of surface. We therefore make no claim in either direction about the relation between X_8 and X_{III} , beyond the shared trace formula – which follows in any case from the shared transcendental lattice. We do not claim the modular form itself as new anywhere in this paper.

Outline. Section 2 defines $\Sigma_{\text{III}}(p)$, reconstructs its hypercuboid origin, and proves the projective reduction to $T(p)$. Section 3 proves the elementary $p \equiv 3 \pmod{4}$ vanishing. Section 4 constructs the elliptic K3 surface X_{III} , establishes its three I_2^* fibers, and proves full rational 2-torsion and torsion completeness. Section 5 proves that every fiber component of X_{III} is defined over $\mathbb{Q}$, via the graph-rigidity argument, and derives $\text{Tr}(F_p \mid \text{NS}(X_{\text{III}})) = 20p$, $\text{disc NS}(X_{\text{III}}) = 4$, and $T(X_{\text{III}}) \cong \text{diag}(2, 2)$. Section 6 identifies the associated CM field $\mathbb{Q}(i)$ and states Livné’s modularity theorem. Section 7 reconstructs, from scratch, the finite twist-elimination argument identifying the governing newform as 16.3.c.a, untwisted. Section 8 re-derives the point-count ledger and proves the paper’s main theorem, Theorem 8.6. Section 9 compares the two resolved branches of the hypercuboid classification. Appendices A–F, containing the longer explicit local computations referenced throughout the text, remain to be written.

2. THE CLASS III CHARACTER SUM

2.1. Hypercuboid origin. We recall (from [1], whose notation we follow exactly) the construction sufficiently for this paper to be read independently. A *clean finite-field hypercuboid residue system* on four coordinates $A_1, A_2, A_3, A_4 \in \mathbb{F}_p$ subject to $\sum_i A_i = 0$ is governed by the seven nonzero pairwise-sum forms $L_J(x) = \sum_{j \in J} x_j$, $\emptyset \neq J \subseteq \{1, 2, 3\}$, in the parametrization $A_1 = x_0, A_2 = x_1, A_3 = x_2, A_4 = -(x_0 + x_1 + x_2)$. The natural S_4 -symmetry permuting $A_1, \dots, A_4$ acts on the seven six-element subsets of these seven forms with exactly three orbits, giving the three character-sum classes $\Sigma_I, \Sigma_{II}, \Sigma_{III}$. Classes I and II each omit a form associated to a *singleton*-type element of the permutation action; Class III omits the form associated to a *pair*-type element, $L_3 = x_0 + x_1$ (the form fixed, up to sign, by the double transposition $(13)(24) \in S_4$) – a genuinely different orbit, with no permutation relating Class III to Classes I or II.

Definition 2.1. For an odd prime p ,

$$\Sigma_{\text{III}}(p) := \sum_{x_0, x_1, x_2 \in \mathbb{F}_p} \chi(P_{\text{III}}(x_0, x_1, x_2)), \quad P_{\text{III}}(x_0, x_1, x_2) := x_0 x_1 x_2 (x_0 + x_2)(x_1 + x_2)(x_0 + x_1 + x_2).$$

P_{III} is manifestly homogeneous of degree 6 (a product of six linear forms) and vanishes identically on the hyperplane $x_0 = 0$.

2.2. Projective reduction.

Lemma 2.2 (Projective reduction). *For every odd prime p ,*

$$\Sigma_{\text{III}}(p) = (p-1)T(p), \quad T(p) := \sum_{x, t \in \mathbb{F}_p} \chi(xt(t+1)(x+t)(x+t+1)).$$

Proof. Since P_{III} is homogeneous of degree $6 = 2 \cdot 3$, for $z \in \mathbb{F}_p^\times$ and $v \in \mathbb{F}_p^3$ we have $P_{\text{III}}(zv) = z^6 P_{\text{III}}(v)$, and $z^6 = (z^3)^2$ is a nonzero square; hence $\chi(P_{\text{III}}(zv)) = \chi(P_{\text{III}}(v))$, so $\chi \circ P_{\text{III}}$ is constant on each punctured line through the origin in $\mathbb{F}_p^3$ (*no exceptional projective direction and no hidden*

sign: the degree being even is exactly what removes any dependence on the choice of nonzero scalar z , for every line, uniformly). The origin contributes $\chi(0) = 0$.

Every line not contained in the hyperplane $x_0 = 0$ has a unique representative of the form $(1, x_1, x_2)$, and direct substitution gives

$$P_{\text{III}}(1, x_1, x_2) = 1 \cdot x_1 \cdot x_2 \cdot (1 + x_2)(x_1 + x_2)(1 + x_1 + x_2).$$

Renaming $x_1 = x$, $x_2 = t$ (*not* the reverse assignment — reordering the factors of $x_1 x_2 (1 + x_2)(x_1 + x_2)(1 + x_1 + x_2)$ under $x_1 \mapsto x$, $x_2 \mapsto t$ gives $x t (1 + t)(x + t)(1 + x + t)$, matching the target only with x_2 , not x_1 , playing the role of t),

$$P_{\text{III}}(1, x, t) = x t (t + 1)(x + t)(x + t + 1),$$

the summand of $T(p)$; there are exactly p^2 such lines, one for each $(x, t) \in \mathbb{F}_p^2$, with no omission or repetition (distinct (x, t) give distinct representatives, since the leading coordinate is fixed at 1). Every line contained in $x_0 = 0$ has $P_{\text{III}}(0, x_1, x_2) = 0$ identically (the *zero vector* and, more generally, the entire hyperplane $x_0 = 0$, contributing $\chi(0) = 0$ throughout — correctly accounted for as zero-contributing, not omitted, matching the analogous treatment for Σ_I in [1]). Each punctured line has $p - 1$ nonzero points, all giving the same character value by the homogeneity argument above (the *orbit size* is $p - 1$ uniformly, with no exceptional orbit of different size), so

$$\Sigma_{\text{III}}(p) = \sum_{v \in \mathbb{F}_p^3} \chi(P_{\text{III}}(v)) = (p - 1) \sum_{\text{lines}} \chi(P_{\text{III}}(v_\ell)) = (p - 1)(T(p) + 0) = (p - 1)T(p). \quad \square$$

Remark 2.3. This is the identical projectivization argument used for $\Sigma_I(p)$ in [1], applied fresh to the genuinely different polynomial P_{III} ; we have re-verified every step above (homogeneity, the vanishing on $x_0 = 0$, the orbit count, and the absence of any hidden sign) independently for this paper rather than inferring the reduction by analogy.

3. ELEMENTARY VANISHING FOR $p \equiv 3 \pmod{4}$

Proposition 3.1 (Elementary mod-4 vanishing). $\Sigma_{\text{III}}(p) = 0$ for every prime $p \equiv 3 \pmod{4}$.

Proof. Define the linear map $\tau : \mathbb{F}_p^3 \rightarrow \mathbb{F}_p^3$ by

$$\tau(x_0, x_1, x_2) = (x_2, -(x_0 + x_1 + x_2), x_0).$$

τ is an involution. Writing $\tau(x_0, x_1, x_2) = (y_0, y_1, y_2)$,

$$\tau(y_0, y_1, y_2) = (y_2, -(y_0 + y_1 + y_2), y_0) = (x_0, -(x_2 - x_0 - x_1 - x_2 + x_0), x_2) = (x_0, x_1, x_2),$$

using $y_0 + y_1 + y_2 = x_2 - (x_0 + x_1 + x_2) + x_0 = -x_1$; hence $\tau \circ \tau = \text{id}$, so τ is a bijection of $\mathbb{F}_p^3$ (a linear involution, matrix square equal to the identity).

The sign identity. A direct expansion (using the factorization of P_{III} in Definition 2.1) shows the polynomial identity

$$P_{\text{III}}(\tau(x_0, x_1, x_2)) = -P_{\text{III}}(x_0, x_1, x_2)$$

holds identically in $\mathbb{Z}[x_0, x_1, x_2]$ — not merely at each tested prime, but as an exact cancellation of every monomial. (*We have independently re-verified this identity by direct symbolic expansion; see the reproducibility appendix.*)

Fixed points. Since τ is a linear involution of $\mathbb{F}_p^3$ (p odd, so 2 is invertible), its fixed locus is a linear subspace, given by $x_2 = x_0$ and $x_1 = -(x_0 + x_1 + x_2)$, i.e. $x_1 = -x_0$ (using $x_2 = x_0$). On this locus, $P_{\text{III}}(x_0, -x_0, x_0)$ vanishes identically — a direct consequence of the factor $(x_0 + x_1 + x_2)$ in the definition of P_{III} , since $x_0 + x_1 + x_2 = x_0 - x_0 + x_0 = x_0$ does *not* vanish there in general, so the vanishing on the fixed locus in fact comes from the factor $(x_1 + x_2) = (-x_0 + x_0) = 0$ — and this vanishing is exactly the internal-consistency check required before accepting the sign identity: $P_{\text{III}}(\tau x) = -P_{\text{III}}(x)$ applied at a fixed point $x = \tau x$ would force $P_{\text{III}}(x) = -P_{\text{III}}(x)$, i.e.

$2P_{\text{III}}(x) = 0$, i.e. $P_{\text{III}}(x) = 0$ (as p is odd) – consistent with, and independently confirmed by, the direct vanishing just exhibited.

Conclusion. Relabeling the summation variable in Definition 2.1 by the bijection τ ,

$$\Sigma_{\text{III}}(p) = \sum_x \chi(P_{\text{III}}(x)) = \sum_x \chi(P_{\text{III}}(\tau x)) = \sum_x \chi(-P_{\text{III}}(x)) = \chi(-1) \sum_x \chi(P_{\text{III}}(x)) = \chi(-1) \Sigma_{\text{III}}(p),$$

an identity valid for *every* odd prime p . When $p \equiv 3 \pmod{4}$, $\chi(-1) = -1$, giving $\Sigma_{\text{III}}(p) = -\Sigma_{\text{III}}(p)$, hence $\Sigma_{\text{III}}(p) = 0$. $\square$

Remark 3.2. The identity $\Sigma_{\text{III}}(p) = \chi(-1)\Sigma_{\text{III}}(p)$ is unconditional for every odd p ; it becomes vacuous, not false, when $\chi(-1) = 1$ (i.e. $p \equiv 1 \pmod{4}$), which is exactly where a genuinely different tool – the K3-geometric and modular argument of Sections 4–5 and of this paper’s later sections – becomes necessary. We note briefly, for the reader, that the same $p \equiv 3 \pmod{4}$ vanishing will reappear in this paper’s modularity section as a consequence of the vanishing of $a_p(16.3.c.a)$ at primes inert in $\mathbb{Q}(i)$, and that $\mathbb{Q}(i)$ ’s inert primes are exactly $p \equiv 3 \pmod{4}$. **These are two logically independent proofs of the same numerical fact**, not one argument in two guises: the present proof never refers to the K3 surface, the modular form, or any arithmetic-geometric structure, and the later CM-inertness proof will never refer to the raw sum P_{III} or to τ . The coincidence in residue class is specific to this surface’s discriminant, not a structural necessity (had the discriminant instead been, say, 8, CM inertness would occur on a different residue class while the present argument, making no reference to any surface, would be unaffected). No modularity is used anywhere in this section.

4. THE ELLIPTIC K3 SURFACE

4.1. Affine model and Weierstrass equation. Writing $T(p) = \#V_{\text{III}}(\mathbb{F}_p) - p^2$ for the affine surface

$$V_{\text{III}} : Y^2 = X(X+1)T(T+1)(X+T+1)$$

(the standard character-sum-to-point-count dictionary applied to the summand of $T(p)$, with T here serving as the coordinate on the base of the elliptic fibration constructed below – we retain this name throughout, rather than relabeling to t , to match the notation already fixed by this project’s exploratory record), we complete the cube in X :

$$X(X+1)(X+T+1) = X^3 + (T+2)X^2 + (T+1)X.$$

Multiplying through by $T^2(T+1)^2$ and substituting $X \mapsto T(T+1)X$, $Y \mapsto T(T+1)Y$ clears the leading coefficient $T(T+1)$ on the cubic, giving the minimal Weierstrass model

$$V_{\text{III}} : Y^2 = X^3 + a_2(T)X^2 + a_4(T)X, \quad a_2(T) = T(T+1)(T+2), \quad a_4(T) = T^2(T+1)^3,$$

with $a_6 = 0$.

4.2. Weierstrass invariants and Kodaira fibers. The standard formulas give

$$c_4(T) = 16T^2(T+1)^2(T^2+T+1), \quad c_6(T) = -32T^3(T-1)(T+1)^3(T+2)(2T+1),$$

$$\Delta(T) = 16T^8(T+1)^8, \quad j(T) = c_4(T)^3/\Delta(T).$$

Proposition 4.1 (Three I_2^* fibers). *V_{III} has bad fibers exactly at $T = 0$, $T = -1$, and $T = \infty$, each of Kodaira type I_2^* , and this Weierstrass model is minimal at every place.*

Proof. At $T = 0$: direct series expansion gives $v_T(c_4) = 2$, $v_T(c_6) = 3$, $v_T(\Delta) = 8$ (the last read off directly from the explicit factored form of Δ). Tate’s algorithm’s criterion for Kodaira type I_n^* is $v(c_4) \geq 2$, $v(c_6) \geq 3$, $v(\Delta) = n+6$; here $v(\Delta) = 8$ gives $n = 2$, i.e. type I_2^* . Since $v(c_4) = 2 < 4$, the model cannot be further reduced at $T = 0$ (Kraus’s minimality criterion), so it is already minimal there. The identical computation applies at $T = -1$, by the symmetric factor $(T+1)^8$ in Δ .

At $T = \infty$: $\deg c_4 = 6$, $\deg c_6 = 9$, $\deg \Delta = 16$. Since a K3 elliptic surface's total discriminant degree is 24, the fiber at infinity contributes $v_\infty(\Delta) = 24 - 16 = 8$. Substituting $S = 1/T$ with the weight-matched rescaling $\tilde{X} = S^4 X$, $\tilde{Y} = S^6 Y$ gives $\tilde{a}_2(S) = S^4 a_2(1/S) = S(S+1)(2S+1)$, $\tilde{a}_4(S) = S^8 a_4(1/S) = S^3(S+1)^3$, from which $v_S(\tilde{c}_4) = 2$, $v_S(\tilde{c}_6) = 3$ – again matching the I_2^* criterion with $n = 2$, and again minimal by the same valuation bound.

No other bad fiber exists: $\Delta(T) = 16T^8(T+1)^8$ has no other zero in T , and the three fibers found account for a total Euler-number contribution of $3 \times 8 = 24$, exactly the Euler characteristic of a K3 surface – confirming both that no further singular fiber remains unaccounted for and, combined with the trivial-lattice rank computation of Proposition 4.2 below, that X_{III} (the smooth relatively minimal model of V_{III} over the T -line) is indeed a K3 surface. $\square$

4.3. Trivial lattice and Picard rank.

Proposition 4.2. $\text{Triv}(X_{\text{III}}) = U + D_6 + D_6 + D_6$, of rank $2 + 6 + 6 + 6 = 20$; the Mordell–Weil rank of X_{III} is 0.

Proof. Each I_2^* fiber contributes a root sublattice of type D_6 to the trivial lattice (the reducible-fiber contribution in the Shioda–Tate formula, one summand per bad fiber, of type D_{n+4} for a fiber of type I_n^* [10]); the zero section and general fiber class contribute the hyperbolic plane U . Since X_{III} is a K3 surface, $\rho(X_{\overline{\mathbb{Q}}}) \leq 20$ (the Hodge-theoretic bound $h^{1,1} = 20$); the trivial lattice alone already has rank 20, so $\rho(X_{\overline{\mathbb{Q}}}) = 20$ exactly (a *singular* K3 surface, in the sense of Shioda–Inose), and the Shioda–Tate formula $\text{rank NS}(X_{\text{III}}) = 2 + \sum_v (m_v - 1) + \text{rank MW}$ forces $\text{rank MW} = 0$. $\square$

4.4. Full rational 2-torsion.

Proposition 4.3 (Global factorization).

$$X^3 + a_2(T)X^2 + a_4(T)X = X(X + T(T+1))(X + T(T+1)^2)$$

identically, as polynomials in $\mathbb{Z}[T][X]$.

Proof. Direct expansion of the right-hand side, matched term by term against $a_2(T) = T(T+1)(T+2)$, $a_4(T) = T^2(T+1)^3$ via Vieta's formulas on the roots $0, -T(T+1), -T(T+1)^2$: their sum is $-T(T+1) - T(T+1)^2 = -T(T+1)(1 + (T+1)) = -T(T+1)(T+2) = -a_2(T)$, matching the coefficient of X^2 ; their pairwise-product sum reduces to $T(T+1) \cdot T(T+1)^2 = T^2(T+1)^3 = a_4(T)$, matching the coefficient of X ; the triple product is 0, matching $a_6 = 0$. $\square$

Corollary 4.4 (Full rational 2-torsion). *The three points $P_1 = (0, 0)$, $P_2 = (-T(T+1), 0)$, $P_3 = (-T(T+1)^2, 0)$ are nonzero 2-torsion sections of the elliptic fibration, individually defined over $\mathbb{Q}(T)$, together with the zero section O forming a subgroup of the Mordell–Weil group isomorphic to $(\mathbb{Z}/2)^2$.*

Proof. On a Weierstrass curve $y^2 = x^3 + a_2x^2 + a_4x + a_6$, a point of order exactly 2 satisfies $P = -P = (x, -y)$, forcing $y = 0$; thus every nonzero 2-torsion point has $y = 0$ and x -coordinate a root of the cubic. Proposition 4.3 exhibits three such roots, each a rational function of T with $\mathbb{Z}$ -coefficients, hence each defining a section individually rational over $\mathbb{Q}(T)$. $\square$

Proposition 4.5 (Torsion completeness). $\text{MW}(X_{\text{III}})_{\text{tors}} \cong (\mathbb{Z}/2)^2$ exactly.

Proof. No larger 2-torsion. A cubic polynomial has at most three roots; the three roots exhibited in Proposition 4.3 are pairwise distinct as elements of $\mathbb{Q}(T)$ (their pairwise differences $T(T+1)$, $T(T+1)^2$, $T^2(T+1)$ are nonzero polynomials, vanishing only at $T = 0, -1$, the two finite bad fibers – confirming they are genuinely distinct sections, not accidentally equal). Since every nonzero 2-torsion point has $y = 0$ (as above) and the cubic has no further roots, the full 2-torsion subgroup is exactly $\{O, P_1, P_2, P_3\} \cong (\mathbb{Z}/2)^2$, no larger.

No torsion of order > 2 . By Shioda’s injectivity theorem [10] (see also [11] for the general theory), the specialization map embeds MW_{tors} into the direct sum of component groups $\bigoplus_v \Phi(\text{fiber}_v)$ over the bad fibers. Each of the three I_2^* fibers has component group $\Phi(I_2^*) \cong (\mathbb{Z}/2)^2$ (Kodaira–Néron), an exponent-2 group; a direct sum of exponent-2 groups is itself exponent-2. Hence every element of MW_{tors} has order dividing 2, ruling out torsion of order 4 or higher outright. Combined with the previous paragraph, $\text{MW}_{\text{tors}} = \text{MW}_{\text{tors}}[2] = (\mathbb{Z}/2)^2$ exactly. $\square$

5. RATIONAL 2-TORSION AND NÉRON–SEVERI RATIONALITY

This section contains the paper’s central technical contribution. Its ingredients are standard – see Remark 5.9 for the prior art on the graph-rigidity step – and what is claimed is the route, not the components. X_{III} already has geometric Picard number 20 and Mordell–Weil rank 0 (Proposition 4.2), so unlike the twist-isomorphism argument of [1], no further Mordell–Weil generator is available to help control the arithmetic of $\text{NS}(X_{\text{III}})$. We prove instead that $\text{NS}(X_{\text{III}})$ is defined over $\mathbb{Q}$ directly, without resolving any of the three I_2^* fibers explicitly.

Proposition 5.1 (Rationality of the I_2^* components). *For each of the three I_2^* fibers of X_{III} , every irreducible component is defined over $\mathbb{Q}$.*

We prove Proposition 5.1 in five steps, working first at the representative fiber $T = 0$; the identical argument applies at $T = -1$ and $T = \infty$ (Remark 5.10).

Step 1: multiplicity-one legs.

Lemma 5.2. *The dual graph of a Kodaira fiber of type I_2^* has exactly seven irreducible components: four of multiplicity 1 (“legs,” in two pairs, one pair attached to each end of a length-three path) and three of multiplicity 2 (“spine”). As an abstract graph it is a tree: six edges among seven vertices.*

Proof. This is the extended (affine) Dynkin diagram of type D_6 , the standard dual graph attached to Kodaira type I_2^* [7, 8, 9]. $\square$

Figure 1 depicts this graph, marked as it will be established at $T = 0$ below.

Lemma 5.3 (Fiber-multiplicity verification). *Let π_1 denote the exceptional curve obtained by resolving the singular point of V_{III} ’s minimal model lying over the leg adjacent to the section P_1 at $T = 0$ (constructed by two successive ordinary blow-ups of the total space, detailed in Appendix D). The local equation of the resolved surface, restricted to a transverse slice near this curve and set to $T = 0$, is*

$$y^2 = x$$

up to sign and a unit rescaling – a reduced plane curve, not a perfect square. Consequently T vanishes to order exactly 1 transverse to this component, i.e. it has fiber-multiplicity exactly 1.

Proof sketch. Blowing up the point of V_{III} ’s total space lying over $(X, T, Y) = (0, 0, 0)$ produces an exceptional curve with local equation $y_1^2 - Tx_1(1 + x_1)^2 - \dots = 0$ (higher-order terms in T); this curve has a further singular point at $x_1 = 0$, whose tangent cone is $y_1^2 - Tx_1 = 0$ – a rank-3, nondegenerate quadratic form (an ordinary node). Resolving it by one further blow-up ($x_1 = Tx_4$, $y_1 = Ty_4$) and dividing by the point’s multiplicity T^2 gives a proper transform whose restriction to $T = 0$ is exactly $y_4^2 - x_4 = 0$, as claimed: reduced, smooth, and *not* of the form (polynomial)² that would signal a further multiplicity-2 (spine) component. Full details, including the parallel computation at the remaining leg positions, are given in Appendix D. $\square$

Remark 5.4. We include Lemma 5.3 explicitly, rather than inferring fiber-multiplicity 1 merely from the fact that the relevant singular point resolves in a single blow-up. The two facts are related but distinct: a point could in principle resolve cleanly into a single smooth exceptional curve that is nonetheless a multiplicity-2 (“spine”) component of the fiber divisor, if T happened to vanish to order 2 transverse to it. Lemma 5.3 rules this out by direct computation rather than by assumption.

Step 2: four rational sections, four distinct legs.

Lemma 5.5. *Any section σ of the elliptic fibration meets a fiber F in exactly one irreducible component, necessarily of multiplicity 1.*

Proof. A section meets every fiber transversally in a single point, so $\sigma \cdot F = 1$. Writing $F = \sum_i m_i C_i$ and using that σ meets only one component C_j (with local intersection multiplicity 1, and $\sigma \cdot C_i = 0$ for $i \neq j$),

$$1 = \sigma \cdot F = \sum_i m_i (\sigma \cdot C_i) = m_j \cdot 1,$$

forcing $m_j = 1$. □

Lemma 5.6 (Four distinct legs at $T = 0$). *The zero section O and the sections P_1, P_2, P_3 of Corollary 4.4 meet four pairwise distinct multiplicity-one components of the I_2^* fiber at $T = 0$.*

Proof. By Lemma 5.5, each of O, P_1, P_2, P_3 meets exactly one multiplicity-one component. O meets the fiber's point at infinity of the Weierstrass model, a point never touched by any of the blow-ups resolving the finite singular point of V_{III} at $T = 0$, hence on a component distinct from those met by P_1, P_2, P_3 . Writing $x_1 = X/T$ for the local coordinate used in Lemma 5.3: $P_1 = (0, 0)$ gives $x_1(P_1) = 0$ identically in T , landing (after resolution) on the leg constructed in Lemma 5.3. The sections P_2, P_3 have $x_1(P_2) = -(T + 1)$, $x_1(P_3) = -(T + 1)^2$, both equal to -1 at $T = 0$ – landing in a *different* branch of the resolution, requiring one further level of blow-up (carried out explicitly in Appendix C, Lemma C.2), after which the two further local coordinates evaluate to two *distinct* constants (-1 and -2 , computed exactly, not merely to leading order), placing P_2, P_3 on two further, pairwise distinct, multiplicity-one components – multiplicity verified explicitly for both, not inferred, exactly as for P_1 's leg (Appendix C, Lemma C.2). All four components – O 's, P_1 's, P_2 's, P_3 's – are thus pairwise distinct. □

Step 3: Galois fixes the legs. Since O, P_1, P_2, P_3 are each individually rational over $\mathbb{Q}(T)$ (in fact, each is given by an explicit rational function of T with integer coefficients), each specializes, at the $\mathbb{Q}$ -rational base point $T = 0$, to an individually $\mathbb{Q}$ -rational point of the fiber. By Lemma 5.6, these four points lie on four pairwise distinct multiplicity-one components. Since $\text{Gal}(\overline{\mathbb{Q}}/\mathbb{Q})$ acts on the fiber's components as a permutation preserving multiplicities and the incidence (intersection) graph – the fiber, defined over $\mathbb{Q}$, is Galois-stable as a set of components with an intrinsic, Galois-equivariant incidence structure – and fixes each of the four marked points individually (each being $\mathbb{Q}$ -rational), it must fix each of the four components carrying a marked point individually: moving a component would move its distinguished rational point, contradicting that the point is fixed.

Step 4: graph rigidity. By Lemma 5.2, label the seven components of the I_2^* fiber L_1, L_2 (legs at one end), S_1, S_2, S_3 (spine, a path with S_1 adjacent to L_1, L_2 and S_3 adjacent to L_3, L_4), L_3, L_4 (legs at the other end).

Lemma 5.7 (Graph rigidity). *Let φ be an automorphism of this graph preserving multiplicities and fixing L_1, L_2, L_3, L_4 individually. Then φ is the identity.*

Proof. The graph is a tree (seven vertices, six edges). In a tree, each leaf has a unique neighbor; since L_1 is a leaf with unique neighbor S_1 , and $\varphi(L_1) = L_1$, we get $\varphi(S_1) \in N(\varphi(L_1)) = N(L_1) = \{S_1\}$, so $\varphi(S_1) = S_1$. The identical argument (using L_3) gives $\varphi(S_3) = S_3$. Now S_2 is the *unique* vertex of the graph adjacent to both S_1 and S_3 (immediate from the path structure: no leg is adjacent to both spine ends, and S_2 is the only spine vertex other than S_1, S_3 themselves). Since $\varphi(S_1) = S_1$ and $\varphi(S_3) = S_3$, and φ preserves adjacency, $\varphi(S_2)$ is also a common neighbor of S_1 and S_3 ; by uniqueness, $\varphi(S_2) = S_2$. Thus φ fixes all seven vertices. □

Remark 5.8. The full (unmarked) automorphism group of this graph, respecting multiplicities, has order 8: generated by the two within-end leg transpositions $(L_1 L_2)$, $(L_3 L_4)$, and the wholesale end-swap exchanging $\{L_1, L_2, S_1\}$ with $\{L_3, L_4, S_3\}$ (fixing S_2). We have independently confirmed this by exhaustive enumeration of all $7! = 5040$ permutations of the vertex set, filtered for multiplicity- and edge-preservation, and separately confirmed that exactly one of these eight automorphisms – the identity – fixes all four legs individually (script in Appendix F). **This computation is cited here as an independent check on Lemma 5.7, not as its proof:** the proof above is the short, self-contained graph-theoretic argument, using only the tree structure and uniqueness of adjacency, that the paper relies on.

Remark 5.9 (Prior art for Lemma 5.7). Lemma 5.7 is a *standard corollary*, and we claim no more for it than the packaging. The governing dichotomy is published. For a fiber of type I_n^* , Wazir [5] records that the number of rational components is either $n + 5$ or $n + 3$, determined by the splitting of a quadratic polynomial; Ulmer [6] states the same dichotomy in split/non-split language, with the non-split case leaving exactly one conjugate pair. So at most one pair of components is ever exchanged, and a pair exchanged by a multiplicity-preserving automorphism of the dual graph must be a pair of multiplicity-one legs at a common end. Lemma 5.7 follows at once. Two cautions, both real: the case I_0^* must be treated separately, because its dual graph is $\tilde{D}_4$, whose automorphism group is S_4 rather than the order-8 group above, so the split/non-split dichotomy does not exhaust its monodromy; and the statement presumes a perfect residue field, since over a non-perfect residue field the geometric special fiber of the minimal regular model need not be a Kodaira type at all. Both hypotheses hold throughout this paper, where $n = 2$ and the residue fields are finite or $\mathbb{Q}$. What the lemma is *not* is the assertion that full rational 2-torsion rigidifies every reducible fiber: that is false, as I_4 and I_6 show, since $(\mathbb{Z}/2)^2$ cannot inject into a cyclic component group.

Proof of Proposition 5.1, at $T = 0$. By Step 3, L_1, L_2, L_3, L_4 (in the labeling of Step 4, matched to the four components of Lemma 5.6) are individually Galois-fixed. By Lemma 5.7, the only automorphism of the fiber’s dual graph compatible with this marking is the identity; since Galois acts on the fiber’s components through some such automorphism, it must act as the identity on all seven components, in particular fixing the three spine components S_1, S_2, S_3 . Every component of the $T = 0$ fiber is therefore defined over $\mathbb{Q}$. $\square$

Remark 5.10 (Transfer to $T = -1, \infty$). The argument at $T = -1$ and $T = \infty$ is *not* obtained by transporting the $T = 0$ computation through a symmetry of the surface – no coordinate automorphism of V_{III} carries $T = 0$ literally to $T = -1$ or $T = \infty$ (the naive candidate, $T \mapsto -1 - T$ together with $X \mapsto -X - 1$, sends the defining equation to its own quadratic twist by -1 , not to itself). Consequently the local computation is carried out *separately and explicitly* at each of the two remaining fibers, using the same local parameter substitutions as at $T = 0$ but with each fiber’s own specific coefficients (Appendix C, Lemma C.2): at $T = -1$, the sections specialize to $x_1(P_2) = 1$ (single blow-up) and $x_3(P_1) = 0$, $x_3(P_3) = 1$ (second level); at $T = \infty$, to $x_1(P_3) = -1$ (single blow-up) and $x_3(P_1) = 0$, $x_3(P_2) = -1$ (second level) – nine tangent cones in total across the three fibers, each independently verified rank-3 nondegenerate (Appendix C). What is common to all three fibers is the general shape of the computation (a rank-3 nondegenerate tangent cone of the form $y^2 - c \cdot (\text{base}) \cdot (\text{transverse})$ always resolves to a reduced, multiplicity-one component, proved once as Lemma C.1 and then checked, not assumed, at each of the nine points) and the purely abstract graph-theoretic fact Lemma 5.7, which depends only on the tree structure of the dual graph and not on which fiber it is attached to. Hence Proposition 5.1 holds at all three I_2^* fibers.

Step 5: Néron–Severi basis and Frobenius trace.

Theorem 5.11. $\text{NS}(X_{\text{III}})$ is generated over $\mathbb{Z}$ by the zero section O , the general fiber class F , and the six non-identity components of each of the three I_2^* fibers, a total of $2 + 3 \times 6 = 20$ classes, each

individually defined over $\mathbb{Q}$. Consequently

$$\boxed{\mathrm{Tr}(F_p \mid \mathrm{NS}(X_{\mathrm{III}})) = 20p}$$

for every good odd prime p .

Proof. The generating set and its count of 20 follow from Proposition 4.2 ($\mathrm{Triv}(X_{\mathrm{III}}) = U + D_6 + D_6 + D_6$, of rank $20 = \rho(X_{\overline{\mathbb{Q}}})$, so $\mathrm{NS}(X_{\mathrm{III}}) = \mathrm{Triv}(X_{\mathrm{III}})$ as lattices, with the torsion subgroup of Proposition 4.5 affecting only the index $[\mathrm{NS} : \mathrm{Triv}^{\mathrm{sat}}]$ via $|\mathrm{disc} \mathrm{NS}| = |\mathrm{disc} \mathrm{Triv}|/|\mathrm{MW}_{\mathrm{tors}}|^2$, not the rank: torsion sections are themselves $\mathbb{Z}$ -linear combinations of the already-counted trivial-lattice generators, not independent generators). Each of the 20 generators is individually defined over $\mathbb{Q}$: O, F trivially (global structure of the Weierstrass model, no local subtlety), and the 18 non-identity fiber components by Proposition 5.1 (applied at all three fibers, Remark 5.10). A Frobenius-fixed algebraic (Tate) class contributes eigenvalue exactly p to Frobenius acting on $H_{\mathrm{et}}^2(X_{\mathrm{III}}, \mathbb{Q}_\ell)$; summing over the 20 fixed generators gives $\mathrm{Tr}(F_p \mid \mathrm{NS}(X_{\mathrm{III}})) = 20p$. $\square$




FIGURE 1. The affine- D_6 dual graph of an I_2^* fiber of X_{III} , marked at the representative fiber $T = 0$: legs $L_1, \dots, L_4$ have multiplicity 1, spine components S_1, S_2, S_3 have multiplicity 2. The zero section O and the three nonzero 2-torsion sections P_1, P_2, P_3 occupy the four legs individually (Lemma 5.6), forcing (Lemma 5.7) every component – legs and spine alike – to be individually Galois-fixed.

Lattice consequences. Combining Theorem 5.11 with the lattice structure of Proposition 4.2 and the torsion order of Proposition 4.5: $\mathrm{disc}(\mathrm{Triv}(X_{\mathrm{III}})) = \mathrm{disc}(U) \cdot \mathrm{disc}(D_6)^3 = (-1) \cdot 4^3 = -64$, and $|\mathrm{disc} \mathrm{NS}(X_{\mathrm{III}})| = |\mathrm{disc} \mathrm{Triv}(X_{\mathrm{III}})|/|\mathrm{MW}_{\mathrm{tors}}|^2 = 64/16 = 4$. Since X_{III} is a singular K3 surface, $T(X_{\mathrm{III}}) = \mathrm{NS}(X_{\mathrm{III}})^\perp$ has rank $22 - 20 = 2$ and discriminant 4. A rank-2 even lattice with Gram matrix $\begin{pmatrix} 2a & b \\ b & 2c \end{pmatrix}$ (the even condition forcing both diagonal entries even) has determinant $4ac - b^2$; matching $4ac - b^2 = 4$ with the trivial solution $a = c = 1, b = 0$ gives the Gram matrix $\mathrm{diag}(2, 2)$ directly, so

$$T(X_{\mathrm{III}}) \cong \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix}.$$

This computation uses only the lattice structure established above, at no point invoking any modular form. The identification of this lattice with a specific CM field – via the associated *primitive* binary quadratic form $Q(x, y) = ax^2 + bxy + cy^2$ (note: a, c here are *half* the Gram matrix’s diagonal entries, since the intersection form $v \cdot v = 2ax^2 + 2bxy + 2cy^2$ is twice Q) – and the full modular consequences of this discriminant are the subject of Section 6.

6. TRANSCENDENTAL LATTICE AND COMPLEX MULTIPLICATION

We now separate four distinct notions that are easy to conflate: the *geometric Picard rank*, the *transcendental lattice* as an abstract lattice, the *CM field* attached to that lattice, and the *arithmetic field of definition* of the surface and its Néron–Severi group. Only the last of these required any arithmetic (Galois-theoretic) argument; the first three are purely geometric or lattice-theoretic.

- **Geometric Picard rank.** $\rho(X_{\overline{\mathbb{Q}}}) = 20$ (Proposition 4.2), a statement about X_{III} base-changed to $\overline{\mathbb{Q}}$, with no reference to $\mathbb{Q}$ -rationality of anything.
- **Transcendental lattice.** $T(X_{\text{III}}) \cong \text{diag}(2, 2)$ (end of Section 5), the isomorphism type of a rank-2 even lattice – again a statement with no arithmetic content, following from $T(X_{\text{III}})$ ’s rank and discriminant alone via the classification of rank-2 even lattices (elementary; see below).
- **CM field.** $\mathbb{Q}(i)$, read off from the isomorphism type of $T(X_{\text{III}})$ via the classical dictionary between rank-2 even lattices and imaginary quadratic orders (Shioda–Inose [12]) – again purely lattice-theoretic, requiring no arithmetic input about Galois representations.
- **Arithmetic field of definition.** X_{III} is defined over $\mathbb{Q}$ (its Weierstrass equation has $\mathbb{Z}$ -coefficients), and $\text{NS}(X_{\text{III}})$ is defined over $\mathbb{Q}$ (Theorem 5.11) – these *are* genuinely arithmetic statements. By contrast, the transcendental lattice $T(X_{\text{III}})$, as an abstract $\mathbb{Z}$ -module, has no arithmetic content on its own; what is arithmetic is the *Galois representation* on $T(X_{\text{III}}) \otimes \mathbb{Q}_\ell$, which is the subject of Livné’s theorem below and is *not* determined by the isomorphism type of $T(X_{\text{III}})$ alone.

6.1. From the lattice to the CM field. A rank-2 even lattice with Gram matrix $\begin{pmatrix} 2a & b \\ b & 2c \end{pmatrix}$ has associated *primitive positive-definite binary quadratic form* $Q(x, y) = ax^2 + bxy + cy^2$ (so that the intersection form $v \cdot v$, for $v = xe_1 + ye_2$, equals $2Q(x, y)$), of classical discriminant $D = b^2 - 4ac < 0$. Two such lattices are isomorphic if and only if their associated forms are $\text{SL}_2(\mathbb{Z})$ -equivalent (equivalently, represent the same class in the class group of the order of discriminant D). For $T(X_{\text{III}}) \cong \text{diag}(2, 2)$: $2a = 2, 2c = 2, b = 0$ gives $(a, b, c) = (1, 0, 1)$, $D = 0 - 4 = -4$. Since $|b| \leq a \leq c$ holds ($0 \leq 1 \leq 1$), $(1, 0, 1)$ is already *reduced*, and since the class number $h(-4) = 1$ (the unique reduced form of discriminant -4), this is *the* isomorphism class of discriminant -4 , with associated field

$$\mathbb{Q}(\sqrt{D}) = \mathbb{Q}(\sqrt{-4}) = \mathbb{Q}(2i) = \mathbb{Q}(i).$$

As a sanity check against [1]: their lattice $T(X) \cong \text{diag}(2, 4)$ gives $(a, b, c) = (1, 0, 2)$, $D = -8$, field $\mathbb{Q}(\sqrt{-8}) = \mathbb{Q}(\sqrt{-2})$ – matching their stated discriminant-8, $\mathbb{Q}(\sqrt{-2})$ result exactly, confirming the convention used here is consistent with the companion paper’s own.

6.2. Livné’s theorem: the arithmetic realization.

Theorem 6.1 (Livné [13]). *Let X be a singular K3 surface defined over $\mathbb{Q}$, with transcendental lattice $T(X)$ of discriminant $-D$ ($D > 0$). Then there exists a normalized weight-3 newform f with complex multiplication by $\mathbb{Q}(\sqrt{-D})$ (or an order therein), of level determined by the conductor of the 2-dimensional Galois representation on $T(X) \otimes \mathbb{Q}_\ell$, such that*

$$\text{Tr}(\text{Frob}_p \mid T(X)) = a_p(f)$$

for almost all primes p of good reduction.

Hypothesis check for X_{III} : X_{III} is a singular K3 surface defined over $\mathbb{Q}$ (Propositions 4.1, 4.2, all over $\mathbb{Q}$, not merely $\overline{\mathbb{Q}}$); $T(X_{\text{III}})$ has discriminant -4 (established above). All hypotheses are met, and Theorem 6.1 guarantees the *existence* of some governing newform f , weight 3, CM by $\mathbb{Q}(i)$, at some 2-power level – but does not by itself identify f or rule out a quadratic twist; that is the content of Section 7.

6.3. Identifying the candidate newform. Weight-3 newforms with CM by $\mathbb{Q}(i)$ and rational Hecke eigenvalues are classified by their level; a search of the relevant space at 2-power level identifies the newform

$$f(z) = \eta^6(4z), \quad \text{LMFDB label } 16.3.c.a, \quad f = q - 6q^5 + 9q^9 + 10q^{13} - 30q^{17} + 11q^{25} + \dots$$

as the unique candidate of level 16 with these properties, and confirms its Nebentypus is χ_{-1} , consistent with a weight-3, CM-by- $\mathbb{Q}(i)$ form at level 16 (the Nebentypus of such a form must be the quadratic character of the CM field itself, ramified exactly at the primes dividing the discriminant, here 2).

Remark 6.2 (Known vs. proved). The modular form $f = \eta^6(4z) = 16.3.c.a$ is **not new**: it is exactly Ahlgren, Ono, and Penniston’s own $A(z)$ [2], identified in their Theorem 1.2 as the newform governing their own K3 surface X_8 at their parameter $\lambda = 8$. **What is known in the literature**: this modular form’s existence, its q -expansion, its occurrence as the trace form for X_8 , and (independently, via Huber–Liu–McLaughlin–Ye–Yuan–Zhang [14]) its occurrence for the quartic Fermat K3 surface. To these must be added a fourth known occurrence, established in Section 9.1: X_{III} is isomorphic over $\mathbb{Q}$ to the $\lambda = -1$ member of [2]’s family, whose transcendental lattice $\langle 2 \rangle \oplus \langle 2 \rangle$ and CM by $\mathbb{Q}(i)$ are recorded in [3], citing [4] and [12]. **What is proved in this manuscript** is accordingly not the existence of a discriminant-4 singular K3 with this trace form, which is known, but that the Class III hypercuboid character sum reaches one: that Σ_{III} reduces to $T(p)$, that the attached surface has NS defined over $\mathbb{Q}$ by the argument of Section 5, and that the resulting trace formula is untwisted (Section 7), yielding the exact evaluation of $\Sigma_{\text{III}}(p)$. We do not claim the modular form itself, or the surface, or its lattice, as new anywhere in this paper.

6.4. Bad primes of X_{III} . The level of f (and hence, via Section 7, the set of primes at which any twist of f can be ramified) is constrained by the bad primes of X_{III} as an arithmetic surface over $\mathbb{Z}$. The Weierstrass coefficients $a_2(T) = T(T+1)(T+2)$, $a_4(T) = T^2(T+1)^3$ have $\mathbb{Z}$ -coefficients, and the discriminant $\Delta(T) = 16T^8(T+1)^8$ factors, as a polynomial in T , with no prime dividing every coefficient except through the constant $16 = 2^4$. Consequently, reducing the entire family modulo any odd prime ℓ gives a family with the identical qualitative structure (three I_2^* fibers at the same three points $T = 0, -1, \infty$, by the same valuation computations of Proposition 4.1 carried out over $\mathbb{F}_\ell$ instead of $\mathbb{Q}$) – i.e. *good reduction* of the surface at every odd prime. At $\ell = 2$, the constant factor 16 introduces additional degeneracy not present at other primes. Hence X_{III} **has bad reduction only at $p = 2$** .

7. MODULARITY AND TWIST ELIMINATION

7.1. The candidate set.

Lemma 7.1 (Twist candidate set). *Any quadratic twist of f compatible with X_{III} ’s bad-reduction locus $\{2\}$ (Section 6.4) is by a Dirichlet character of conductor dividing 8.*

Proof. A quadratic twist of a newform unramified outside $\{2\}$ by a character χ_d (Kronecker symbol of fundamental discriminant d) remains unramified outside $\{2\}$ only if d itself is supported only at 2. The fundamental discriminants supported only at the prime 2 are exactly 1, -4 , 8 , -8 , giving the four characters $\{1, \chi(-1), \chi(2), \chi(-2)\}$. $\square$

This candidate set is *forced*, not assumed: it follows directly from X_{III} ’s bad-reduction locus established in Section 6.4, by an elementary ramification argument, before any numerical data is consulted.

7.2. Degeneracy collapse via CM. Since f has CM by $\mathbb{Q}(i)$, its Hecke eigenvalues vanish at primes inert in $\mathbb{Q}(i)$: $a_p(f) = 0$ for $p \equiv 3 \pmod{4}$ (a standard fact for CM newforms – the associated Galois representation is induced from a Hecke character of $\mathbb{Q}(i)$, and its trace at an element not fixing either of the two cosets vanishes). At such primes, *every* candidate twist predicts 0, carrying no discriminating information. Write $F_0(p) := a_p(f)$ (untwisted) and consider the remaining three candidates at $p \equiv 1 \pmod{4}$ (the only primes where discrimination is possible at all): $\chi(-1)(p) = 1$ there identically (definition of $\chi(-1)$), so the twist-by- $\chi(-1)$ candidate equals $F_0(p)$ there, collapsing $\{1, \chi(-1)\}$ into the single testable hypothesis F_0 . Likewise $\chi(-2)(p) = \chi(-1)(p)\chi(2)(p) = \chi(2)(p)$ at $p \equiv 1 \pmod{4}$, collapsing $\{\chi(2), \chi(-2)\}$ into the single testable hypothesis $F_1(p) := \chi(2)(p)a_p(f)$.

Lemma 7.2 (Exhaustiveness). *After collapse, exactly two candidates – F_0 and F_1 – remain to be distinguished, and they exhaust the four-element candidate set of Lemma 7.1 at every prime where any distinguishing test is possible.*

Proof. Immediate from the two collapses above: at $p \equiv 3 \pmod{4}$ all four candidates coincide (all 0); at $p \equiv 1 \pmod{4}$ all four candidates reduce, by the identities just verified, to exactly F_0 or F_1 . $\square$

7.3. The discriminating computation at $p = 5$.

Proposition 7.3 (Twist elimination). *F_1 is not the correct trace formula for $T(X_{\text{III}})$; F_0 survives.*

Proof. By Lemma 7.2, it suffices to exhibit one prime $p \equiv 1 \pmod{4}$ where $F_0(p) \neq F_1(p)$ and to compute $T(p)$ there directly. Take $p = 5$: $5 \equiv 1 \pmod{4}$ (on the support), and $\chi(2)(5) = -1$ (since $5 \not\equiv \pm 1 \pmod{8}$), so $F_1(5) = -F_0(5)$; as $a_5(f) = -6 \neq 0$, $F_0(5) = -6 \neq F_1(5) = 6$, genuinely different predictions. Direct computation of the finite sum $T(5) = \sum_{x,t \in \mathbb{F}_5} \chi(xt(t+1)(x+t)(x+t+1))$ (a sum of 25 terms, exactly evaluable) gives $T(5) = -6 = F_0(5) \neq F_1(5)$. Hence F_1 is refuted at $p = 5$, and by Lemma 7.2 (only F_0, F_1 possible), F_0 is the correct formula. $\square$

Remark 7.4. **This is not a claim that F_0 “matches at $p = 5$.”** The logical structure is: Lemma 7.1 proves the candidate set is finite and exhausted by $\{1, \chi(-1), \chi(2), \chi(-2)\}$, *before* any numerical value is consulted; Lemma 7.2 proves this finite set collapses to exactly two testable hypotheses; and Proposition 7.3 then uses a single exact value to *eliminate* the one candidate that could still have been correct, leaving only F_0 as a matter of logical elimination over an already-proved-exhaustive set – not a pattern observed to hold at one tested point among many untested ones.

Theorem 7.5 (Exact transcendental trace). *For every odd prime $p \neq 2$,*

$$\boxed{\text{Tr}(F_p \mid T(X_{\text{III}})) = a_p(\mathbf{16.3.c.a})}.$$

Explicitly, by cases: for $p \equiv 1 \pmod{4}$ (p split in $\mathbb{Q}(i)$), this is Proposition 7.3; for $p \equiv 3 \pmod{4}$ (p inert in $\mathbb{Q}(i)$), $a_p(f) = 0$ (CM vanishing at inert primes, independent of the twist question) and separately $\text{Tr}(F_p \mid T(X_{\text{III}})) = 0$ follows from Livné’s theorem applied to any of the (here indistinguishable) twist candidates, all of which agree at 0.

8. EXACT EVALUATION

We now derive the point-count ledger relating $T(p)$ to X_{III} ’s point count *from scratch*, independently of any earlier stage of this project’s research record, and combine it with Theorem 5.11 and Theorem 7.5 via the Lefschetz trace formula.

8.1. The ledger, re-derived. Write $\#X_{\text{III}}(\mathbb{F}_p)$ for the point count of the smooth relatively minimal model, fibered over $\mathbb{P}_T^1(\mathbb{F}_p)$ ($p+1$ points: $p-2$ good finite T , and the three bad fibers $T=0, -1, \infty$).

Lemma 8.1 (Affine bad-fiber count). *At $T=0$ and at $T=-1$, the affine curve $Y^2 = X(X+1)T(T+1)(X+T+1)$ has exactly p points over $\mathbb{F}_p$.*

Proof. At $T=0$ or $T=-1$, the factor $T(T+1)$ vanishes, so the right-hand side is identically 0 for every X , giving $Y=0$ for each of the p values of X : exactly p points. $\square$

Lemma 8.2 (Good-fiber affine count). *For each good finite T , the affine curve $Y^2 = X(X+1)T(T+1)(X+T+1)$ has exactly $p - a_T$ points, where a_T is the Frobenius trace of the elliptic curve E_T obtained from this cubic.*

Proof. For $T(T+1) \neq 0$, the rescaling $X' = T(T+1)X$, $Y' = T(T+1)Y$ used in Section 4 to reach the minimal Weierstrass model is a bijection on $\mathbb{F}_p$ -points (multiplication by the nonzero scalar $T(T+1)$), so the affine point count on the original cubic equals the affine point count on the Weierstrass model, which is $p - a_T$ by the standard formula (projective count $p+1 - a_T$, minus the single point at infinity). $\square$

Proposition 8.3 (Ledger).

$$\#X_{\text{III}}(\mathbb{F}_p) = 1 + p^2 + 20p + T(p).$$

Proof. By definition (Section 4), $\#V_{\text{III}}(\mathbb{F}_p) = p^2 + T(p)$. By Lemmas 8.1–8.2,

$$\#V_{\text{III}}(\mathbb{F}_p) = \sum_{\text{good finite } T} (p - a_T) + p + p = (p-2)p - \sum_{\text{good finite } T} a_T + 2p,$$

so $\sum_{\text{good finite } T} a_T = p^2 - \#V_{\text{III}}(\mathbb{F}_p) = -T(p)$.

On the smooth projective model, each good finite fiber contributes $p+1 - a_T$ points, and each of the three I_2^* fibers – now that Proposition 5.1 has established *every* component of each is individually $\mathbb{F}_p$ -rational – contributes $7(p+1) - 6 = 7p+1$ points (seven rational $\mathbb{P}^1$'s, each with $p+1$ points, glued at six nodes each counted twice). Hence

$$\begin{aligned} \#X_{\text{III}}(\mathbb{F}_p) &= \sum_{\text{good finite } T} (p+1 - a_T) + 3(7p+1) \\ &= (p-2)(p+1) - \sum_{\text{good finite } T} a_T + 21p + 3 \\ &= (p^2 - p - 2) + T(p) + 21p + 3 \quad (\text{using } \sum a_T = -T(p) \text{ above}) \\ &= 1 + p^2 + 20p + T(p). \end{aligned} \quad \square$$

Remark 8.4. Every constant appearing here – the p -point counts at the two finite bad affine fibers (Lemma 8.1), the $7p+1$ count per resolved I_2^* fiber (a direct consequence of Section 5's now-proved component rationality, not assumed), and the resulting $20p+1$ net correction – has been independently re-derived in this section, not carried over from any earlier stage of this project's research record. (We have additionally cross-checked Proposition 8.3 by direct brute-force computation of $\#V_{\text{III}}(\mathbb{F}_p)$ for several primes, confirming the formula numerically; see the reproducibility appendix. This check is illustrative, not part of the proof.)

8.2. The main theorem.

Theorem 8.5 (Exact evaluation of $T(p)$). *For every odd prime p ,*

$$T(p) = \begin{cases} a_p(16.3.c.a) & p \equiv 1 \pmod{4} \\ 0 & p \equiv 3 \pmod{4}. \end{cases}$$

Proof. The Lefschetz trace formula for a K3 surface gives

$$\#X_{\text{III}}(\mathbb{F}_p) = 1 + p^2 + \text{Tr}(F_p \mid \text{NS}(X_{\text{III}})) + \text{Tr}(F_p \mid T(X_{\text{III}})) = 1 + p^2 + 20p + a_p(f)$$

using Theorem 5.11 and Theorem 7.5. Equating with Proposition 8.3's independent computation of the same quantity, $1 + p^2 + 20p + T(p) = 1 + p^2 + 20p + a_p(f)$, gives $T(p) = a_p(f)$ for every odd p – in particular for $p \equiv 1 \pmod{4}$. For $p \equiv 3 \pmod{4}$, this is exactly Proposition 3.1, already proved without any of the above machinery. $\square$

Theorem 8.6 (Main theorem). *For every odd prime p ,*

$$\Sigma_{\text{III}}(p) = \begin{cases} (p-1) a_p(16.3.c.a) & p \equiv 1 \pmod{4} \\ 0 & p \equiv 3 \pmod{4}. \end{cases}$$

Proof. Immediate from Lemma 2.2 ($\Sigma_{\text{III}}(p) = (p-1)T(p)$) and Theorem 8.5. $\square$

Remark 8.7 (Consistency with the elementary vanishing). Theorem 8.6's $p \equiv 3 \pmod{4}$ case was already proved, unconditionally and without any arithmetic geometry, in Proposition 3.1. The modular route above reproduces the same conclusion by an entirely different mechanism (CM vanishing at inert primes, Remark 3.2) – a nontrivial consistency check: had the two computations disagreed, something in Sections 4–8 would have been wrong. They agree. We do not introduce a further mod-8 (or finer) case split: no genuinely new phenomenon appears at that finer level, and the two-case statement above is already the sharpest form the evidence supports.

9. COMPARISON WITH CLASSES I/II

The following table records the established invariants of both branches of the hypercuboid classification's resolved sectors, side by side.

	Classes I/II [1]	Class III (this paper)
$ \text{disc } T(X) $	8	4
CM field	$\mathbb{Q}(\sqrt{-2})$	$\mathbb{Q}(i)$
Modular form level	8	16
Mordell–Weil rank	1 (over a quadratic extension)	0
NS mechanism (§5 vs. [1])	MW section + twist	2-torsion + rigidity
Bad fiber count	4	3 ($3 \times I_2^*$)
Twist	$\chi(2)$ -twisted	untwisted

We record this table strictly as an **observation about the two resolved branches of the present hypercuboid system**, not as a theorem. In particular, we do not claim, and this table does not establish, any general correspondence

$$\text{hypercuboid constraint class} \longrightarrow \text{CM field},$$

nor any general principle governing which NS-rationality mechanism (explicit extra section versus torsion-and-rigidity) applies to a given branch. Only two branches of the classification have been resolved to date (Class III's $p \equiv 1 \pmod{4}$ sector being the last one), and two data points do not establish a law. Whether the pattern visible here (distinct discriminants, distinct CM fields, distinct proof mechanisms) reflects something structural about the hypercuboid construction, or is simply how these two particular sums happen to resolve, is an open question we leave for future work, stated here only as a structural observation.

9.1. Precise relation to Ahlgren–Ono–Penniston. We record here, precisely, the outcome of a targeted re-examination of [2] against each of the following six possibilities: that AOP already contains (i) this exact Class III character sum; (ii) this exact hypercuboid reduction; (iii) this exact K3 model X_{III} ; (iv) the graph-rigidity NS-rationality argument of Section 5; (v) merely the same modular form; or (vi) a surface birationally or isomorphically related to X_{III} .

[2] is a paper about their own one-parameter family X_λ , with no combinatorial hypercuboid motivation and no notion of a “Class III” sum; possibilities (i) and (ii) do not apply.

Possibilities (iii) and (vi) *do* apply, and an earlier version of this paper stated the opposite. The comparison made there was against the single parameter $\lambda = 8$, and it did not extend over the family. As recorded in the Introduction, $T(p) = A(-1, p)$ at every odd prime tested, and the branch sextic of X_{III} is carried onto that of the $\lambda = -1$ member by an explicit $\mathbb{Q}$ -linear substitution with no residual scalar, so that $X_{\text{III}} \cong X_{-1}$ over $\mathbb{Q}$. The exact sum and the exact surface of this paper are therefore *not* absent from [2]’s family: they are its $\lambda = -1$ member, reached from an unrelated combinatorial starting point.

Nor is the resulting arithmetic new. The surface X_{-1} is classically studied: van Geemen and Top [3] record its transcendental lattice as $\langle 2 \rangle \oplus \langle 2 \rangle$ of discriminant 4, citing Persson [4], and observe that Shioda and Inose [12] identify it as the desingularised quotient of $E_i \times E_i$, $E_i = \mathbb{C}/\mathbb{Z}[i]$, giving CM by $\mathbb{Q}(i)$. Possibility (v) accordingly also applies, and in a stronger form than the $\lambda = 8$ coincidence: the transcendental lattice, the discriminant and the CM field are all prior art.

What does *not* appear in [2], in [3], or in the literature reviewed is possibility (iv): the argument of Section 5, which establishes NS-rationality from full rational 2-torsion plus rigidity of the fiber’s incidence graph, without resolving any singular fiber. Even here we claim only that the packaging is ours. The graph-rigidity step is a standard corollary of published facts about the Galois action on I_n^* fibers: Wazir [5] records that an I_n^* fiber has either $n + 5$ or $n + 3$ rational components, according to the splitting of a quadratic, and Ulmer [6] gives the same dichotomy in split/non-split form; only a conjugate pair of multiplicity-one components can ever be exchanged, which is precisely Lemma 5.7’s conclusion. See the remark following that lemma.

Summarising the provenance honestly: the character-sum identity, the hypercuboid reduction, the elementary involution proof and the NS-rationality route are this paper’s; the surface, its transcendental lattice, its CM field and its modular form are not. The contribution is an independent derivation and a combinatorial realization of a known singular-K3 arithmetic sector, not the discovery of one.

One qualification cuts the other way and should be recorded with the same care. The parameter reached here is exactly the one at which [2]’s own evaluation degenerates, since their closed form carries a factor $\chi(\lambda + 1)$ and an inverse of $\lambda + 1$, and [3]’s isogeny theorem likewise excludes $t = -1$. The exact evaluation of $\Sigma_{\text{III}}(p)$ obtained here is therefore not a specialization of [2]’s formula; it is obtained by a route that remains valid where that formula does not apply. What is prior art is the surface and its lattice-theoretic and CM data, not the evaluation of the sum at this parameter.

10. DISCUSSION

This paper resolves the Class III sector of the hypercuboid classification left open by [1], completing the finite-field hypercuboid system’s four-coordinate case. The central technical contribution – proving NS-rationality for all three I_2^* fibers of X_{III} via full rational 2-torsion and rigidity of the fiber’s incidence graph, without resolving any fiber explicitly – is a different route from [1]’s twist-isomorphism argument, and may be of use for other singular K3 surfaces exhibiting full rational 2-torsion at their reducible fibers. Its graph-rigidity step is a standard corollary of published I_n^* component behaviour (Remark 5.9); the assembly is what is claimed.

We make no claim to any new general theorem in the theory of K3 surfaces or modular forms: the machinery used here (Kodaira’s classification, Shioda–Tate, Livné’s modularity theorem) is entirely classical, and the contribution is the specific application to the surface that the Class III hypercuboid reduction produces. Nor, as Section 9.1 sets out, is that surface new: it is the $\lambda = -1$ member of [2]’s family, and its transcendental lattice, discriminant and CM field are recorded in [3], [4] and [12]. The hypercuboid route reaches a known singular-K3 arithmetic sector independently; it does not exhibit a new one.

One further caution, drawn from the same family. The branch arrangement $\{x, y, x+1, y+1, x+\lambda y, \text{line at infinity}\}$ has the *same* incidence combinatorics – three triple points and six double points, the same matroid, the same singularity pattern – for every λ outside $\{0, -1, \infty\}$, while its arithmetic varies sharply across that stratum: the generic member has Picard number 19 and no CM at all, and particular members are singular with differing CM data. **Bare incidence combinatorics therefore do not determine the arithmetic sector within this family.** We draw only that negative conclusion. It is not a claim that combinatorics is irrelevant to arithmetic, nor that no combinatorics-to-arithmetic relationship is possible; an enriched invariant carrying realization or moduli data could still be informative, and that question is untouched here. In particular this paper does **not** claim, and its two data points could not support, any general law sending a hypercuboid class to a CM field.

This paper does not address the classical integer perfect-hypercuboid problem, to which the finite-field system studied here and in [1] is only a motivating, logically separate construction.

APPENDIX A. CHARACTER-SUM REDUCTION, IN FULL

This appendix reconstructs Lemma 2.2's reduction with every substitution written out. A specific variable-labeling error is possible at exactly one step below (the affine-slice renaming); writing out both the correct and the incorrect substitution explicitly, and verifying by direct symbolic comparison which one matches the target sum, is meant to make that error impossible to reintroduce unnoticed.

Homogeneity. $P_{\text{III}}(x_0, x_1, x_2) = x_0 x_1 x_2 (x_0 + x_2)(x_1 + x_2)(x_0 + x_1 + x_2)$ is a product of six linear forms, hence homogeneous of degree 6: for $z \in \mathbb{F}_p^\times$,

$$P_{\text{III}}(zx_0, zx_1, zx_2) = z^6 P_{\text{III}}(x_0, x_1, x_2).$$

Since $6 = 2 \cdot 3$ is even, $z^6 = (z^3)^2$ is a square in $\mathbb{F}_p^\times$ for *every* $z \neq 0$. Hence $\chi(P_{\text{III}}(zv)) = \chi(P_{\text{III}}(v))$ for every $z \in \mathbb{F}_p^\times$ and every $v \in \mathbb{F}_p^3$: **no exceptional direction, no hidden sign, valid uniformly over every line through the origin.**

The zero vector and the hyperplane $x_0 = 0$. $P_{\text{III}}(0, x_1, x_2) = 0$ identically (manifest, x_0 a factor), so $\chi(P_{\text{III}}(v)) = 0$ for every v with $v_0 = 0$, including $v = 0$. These are exactly the points *not* covered by the affine-slice parametrization below, and they are correctly accounted for as contributing 0, not silently dropped.

Affine slice and the corrected substitution. Every point v with $v_0 \neq 0$ lies on a unique line through the origin meeting the hyperplane $x_0 = 1$, at the point $(1, v_1/v_0, v_2/v_0)$. Writing this representative as $(1, x_1, x_2)$:

$$P_{\text{III}}(1, x_1, x_2) = 1 \cdot x_1 \cdot x_2 \cdot (1 + x_2) \cdot (x_1 + x_2) \cdot (1 + x_1 + x_2).$$

This is the exact step at which a mislabeling is possible: renaming $x_1 \mapsto t, x_2 \mapsto x$ gives $tx(1+x)(t+x)(1+t+x)$, in which the factor $(1+x)$ does *not* match the target sum's $(t+1)$ – demonstrably wrong, verifiable by direct symbolic comparison. The **correct** renaming is $x_1 \mapsto x, x_2 \mapsto t$:

$$P_{\text{III}}(1, x, t) = x \cdot t \cdot (1+t) \cdot (x+t) \cdot (1+x+t) = xt(t+1)(x+t)(x+t+1),$$

which *is* the summand of $T(p)$. As a symbolic identity (independent of any prime p , checked as an equality of integer polynomials):

$$P_{\text{III}}(1, x, t) - xt(t+1)(x+t)(x+t+1) = 0 \quad \text{in } \mathbb{Z}[x, t].$$

(Verified by direct expansion; see the companion research record's scripts, which re-derive exactly this cancellation to literal zero.)

Orbit size and the factor $p-1$. For fixed $(x, t) \in \mathbb{F}_p^2$, the line through the origin and $(1, x, t)$ has exactly $p-1$ nonzero points (all scalar multiples $z \cdot (1, x, t)$, $z \in \mathbb{F}_p^\times$), and, by homogeneity,

$\chi(P_{\text{III}}(\cdot))$ is constant along it. Distinct pairs (x, t) give distinct lines (the representative $(1, x, t)$ is unique per line meeting $x_0 = 1$), and every line with $v_0 \neq 0$ arises this way, so there is a bijection between $\mathbb{F}_p^2$ and the set of such lines. Hence

$$\Sigma_{\text{III}}(p) = \sum_{v \in \mathbb{F}_p^3} \chi(P_{\text{III}}(v)) = (p-1) \underbrace{\sum_{(x,t) \in \mathbb{F}_p^2} \chi(P_{\text{III}}(1, x, t))}_{\text{lines with } v_0 \neq 0} + \underbrace{0}_{\text{hyperplane } x_0=0} = (p-1)T(p).$$

APPENDIX B. WEIERSTRASS MODEL AND SINGULAR FIBERS, IN FULL

From the affine surface to the Weierstrass model. Starting from $V_{\text{III}} : Y^2 = X(X + 1)T(T + 1)(X + T + 1)$, expand the cubic in X :

$$X(X+1)(X+T+1) = (X^2+X)(X+T+1) = X^3+(T+1)X^2+X^2+(T+1)X = X^3+(T+2)X^2+(T+1)X.$$

So $Y^2 = T(T+1)[X^3+(T+2)X^2+(T+1)X]$. Writing $A := T(T+1)$ and substituting $X' = AX$, $Y' = AY$ (equivalently: multiply both sides by A^2 and regroup):

$$A^2Y^2 = A^3X^3 + A^3(T+2)X^2 + A^2 \cdot A(T+1) \cdot X \implies Y'^2 = X'^3 + A(T+2)X'^2 + A^2(T+1)X',$$

and $A(T+2) = T(T+1)(T+2) = a_2(T)$, $A^2(T+1) = T^2(T+1)^2 \cdot (T+1) = T^2(T+1)^3 = a_4(T)$, matching Section 4 exactly:

$$a_2(T) = T(T+1)(T+2), \quad a_4(T) = T^2(T+1)^3, \quad a_6 = 0.$$

Invariants. With $a_1 = a_3 = a_6 = 0$: $b_2 = 4a_2$, $b_4 = 2a_4$, $b_6 = 0$, $b_8 = 4a_2 \cdot 0 - a_4^2 = -a_4^2$, and

$$c_4 = b_2^2 - 24b_4 = 16a_2^2 - 48a_4, \quad c_6 = -b_2^3 + 36b_2b_4 - 216b_6 = -64a_2^3 + 288a_2a_4,$$

$$\Delta = \frac{c_4^3 - c_6^2}{1728}.$$

Substituting and factoring (direct polynomial arithmetic, independently re-verified by symbolic computation):

$$c_4(T) = 16T^2(T+1)^2(T^2+T+1), \quad c_6(T) = -32T^3(T-1)(T+1)^3(T+2)(2T+1),$$

$$\Delta(T) = 16T^8(T+1)^8, \quad j(T) = c_4(T)^3/\Delta(T).$$

Fiber at $T = 0$: local parameter T itself. Series expansion: $c_4(T) = 16T^2 + O(T^3)$ (leading coefficient from $16 \cdot 1^2 \cdot 1 = 16$), so $v_T(c_4) = 2$; $c_6(T) = 64T^3 + O(T^4)$ ($-32 \cdot (-1) \cdot 1 \cdot 2 \cdot 1 = 64$), so $v_T(c_6) = 3$; $\Delta(T) = 16T^8(1+T)^8$, and $(1+T)^8$ is a unit at $T = 0$, so $v_T(\Delta) = 8$ exactly. Tate's criterion: $v(c_4) \geq 2$, $v(c_6) \geq 3$, $v(\Delta) = n + 6$ gives type I_n^* with $n = 8 - 6 = 2$, i.e. I_2^* . Minimality: $v(c_4) = 2 < 4$, so (Kraus) the model cannot be further reduced – already minimal.

Fiber at $T = -1$: local parameter $s = T+1$. By the manifest symmetry of $\Delta(T) = 16T^8(T+1)^8$ under the factor $(T+1)^8$, the identical computation (with $T \leftrightarrow -1$'s role played by $s = 0$) gives $v_s(c_4) = 2$, $v_s(c_6) = 3$, $v_s(\Delta) = 8$: again I_2^* , again minimal.

Fiber at $T = \infty$: local parameter $S = 1/T$. Since $\deg c_4 = 6$, $\deg c_6 = 9$, $\deg \Delta = 16$, and a K3 elliptic surface's total discriminant degree is 24 (the anticanonical bound), the fiber at infinity has $v_\infty(\Delta) = 24 - \deg \Delta = 24 - 16 = 8$. To see the Kodaira type directly (not merely via the degree count), rescale $\tilde{X} = S^4X$, $\tilde{Y} = S^6Y$ (matching X 's weight 2, Y 's weight 3 against a_2 's degree bound 4, a_4 's bound 8):

$$\tilde{a}_2(S) = S^4a_2(1/S) = S(S+1)(2S+1), \quad \tilde{a}_4(S) = S^8a_4(1/S) = S^3(S+1)^3.$$

Then $\tilde{c}_4(S) = 16\tilde{a}_2^2 - 48\tilde{a}_4$: $\tilde{a}_2^2$ has $v_S = 2$ (leading term S^2), $\tilde{a}_4$ has $v_S = 3$, so the degree-2 term dominates and $v_S(\tilde{c}_4) = 2$. Similarly $\tilde{c}_6 = -64\tilde{a}_2^3 + 288\tilde{a}_2\tilde{a}_4$: $\tilde{a}_2^3$ has $v_S = 3$, $\tilde{a}_2\tilde{a}_4$ has $v_S = 1 + 3 = 4$, so $v_S(\tilde{c}_6) = 3$. Again I_2^* ($n = 8 - 6 = 2$), again minimal ($v(\tilde{c}_4) = 2 < 4$).

No other bad fiber: $\Delta(T) = 16T^8(T+1)^8$ has zeros only at $T = 0, -1$ (finite) and the computed order at $T = \infty$; no other prime factor of T appears.

Euler-number check. Each I_2^* fiber contributes Euler number $v(\Delta) = 8$ (Ogg’s formula, tame case); $3 \times 8 = 24$, exactly the Euler characteristic of a K3 surface – an independent consistency check confirming (i) no bad fiber was missed and (ii) the surface’s total space is indeed K3-type, not merely “some elliptic surface.”

Component/multiplicity pattern. The dual graph of a Kodaira fiber of type I_2^* is the extended Dynkin diagram of type D_6 : seven components, four of multiplicity 1 and three of multiplicity 2 (Lemma 5.2, cited to Kodaira/Néron/Tate; not re-derived from the Weierstrass equation alone in this appendix, but see Appendix C for the explicit local resolution confirming the multiplicity-1 components at each fiber).

APPENDIX C. RATIONAL 2-TORSION AND COMPONENT SPECIALIZATION, IN FULL

The factorization, verified. Matching $X^3 + a_2(T)X^2 + a_4(T)X = X(X + T(T + 1))(X + T(T + 1)^2)$ via Vieta: writing $r_1 = -T(T + 1)$, $r_2 = -T(T + 1)^2$, the sum of the three roots 0, r_1 , r_2 is

$$r_1 + r_2 = -T(T + 1) - T(T + 1)^2 = -T(T + 1)[1 + (T + 1)] = -T(T + 1)(T + 2) = -a_2(T),$$

matching (Vieta) the negative of the X^2 -coefficient; the sum of pairwise products $(0 \cdot r_1 + 0 \cdot r_2 + r_1 r_2)$ is

$$r_1 r_2 = T(T + 1) \cdot T(T + 1)^2 = T^2(T + 1)^3 = a_4(T),$$

matching the X -coefficient; the triple product is 0, matching $a_6 = 0$.

The three sections. $P_1 = (0, 0)$, $P_2 = (-T(T + 1), 0)$, $P_3 = (-T(T + 1)^2, 0)$, each a point on V_{III} ’s Weierstrass model with coordinates that are polynomials in T with integer coefficients – hence each individually defined over $\mathbb{Q}(T)$.

Torsion completeness, fully spelled out. Pairwise differences of the three roots: $r_0 - r_1 = 0 - (-T(T + 1)) = T(T + 1)$; $r_0 - r_2 = T(T + 1)^2$; $r_1 - r_2 = -T(T + 1) + T(T + 1)^2 = T(T + 1)[(T + 1) - 1] = T^2(T + 1)$. All three are nonzero polynomials in $\mathbb{Z}[T]$ (not the zero polynomial), vanishing (as polynomials) only at $T = 0, -1$ – exactly the two finite bad fibers. Hence the three roots are pairwise distinct as elements of $\mathbb{Q}(T)$, so P_1, P_2, P_3 are three genuinely distinct nonzero 2-torsion sections (together with O : four elements). Since any point of order exactly 2 on $y^2 = x^3 + a_2x^2 + a_4x + a_6$ has $y = 0$ (as $P = -P \Leftrightarrow y = -y$), and the cubic has at most three roots, **no further 2-torsion point exists**: the full 2-torsion is exactly $\{O, P_1, P_2, P_3\} \cong (\mathbb{Z}/2)^2$. For torsion of order > 2 : Shioda’s injectivity theorem embeds MW_{tors} into $\bigoplus_v \Phi(\text{fiber}_v)$; each of the three I_2^* fibers has $\Phi(I_2^*) \cong (\mathbb{Z}/2)^2$ (exponent 2), so the direct sum is exponent 2, ruling out any element of order 4 or higher. Hence $\text{MW}_{\text{tors}} \cong (\mathbb{Z}/2)^2$ exactly.

Specialization table. Using $x_1 = X/T$ (local coordinate at $T = 0$), $x_1 = X/s$ ($s = T + 1$, at $T = -1$), and the analogous coordinate at $T = \infty$ (Section 4’s $\tilde{X} = S^4X$ rescaling), together with the second-level blow-up coordinates x_3 used to separate points landing at the same first-level value:

	$T = 0$	$T = -1$	$T = \infty$
O	identity leg	identity leg	identity leg
P_1	$x_1 = 0$ (leg)	$x_3 = 0$ (leg)	$x_3 = 0$ (leg)
P_2	$x_3 = -1$ (leg)	$x_1 = 1$ (leg)	$x_3 = -1$ (leg)
P_3	$x_3 = -2$ (leg)	$x_3 = 1$ (leg)	$x_1 = -1$ (leg)

At each fiber the four values are pairwise distinct (as curves, not merely as numerals – x_1 and x_3 parametrize *different* exceptional components, so a numerical coincidence across columns, e.g. “ -1 ” appearing as both an x_1 -value and an x_3 -value at different fibers, never indicates the same component). This is Lemma 5.6 and Remark 5.10’s content, tabulated explicitly.

Multiplicity-one verification, complete. This subsection proves fiber-multiplicity 1 at *every* marked component in the specialization table above – nine tangent-cone computations in total

(P_1, P_2, P_3 at each of $T = 0, -1, \infty$) – not merely at the single point (P_1 at $T = 0$) treated in the main text’s Lemma 5.3. No component is identified as a leg merely because Lemma 5.2 says four legs must exist somewhere in the fiber: each of the nine points below is checked individually, by exhibiting its own local equation.

Lemma C.1 (General multiplicity criterion). *Suppose, in local coordinates (u, T, y) centered at a singular point of V_{III} ’s total space (with T the base coordinate), the defining equation has the form $F = y^2 - cTu + (\text{terms each carrying an explicit factor } T^2, \text{ or total degree } \geq 4)$, for a nonzero constant c . Then the exceptional curve obtained by blowing up this point in the T -dominant chart ($u = Tu', y = Ty'$) has fiber-multiplicity exactly 1.*

Proof. The tangent cone $y^2 - cTu$ is a rank-3 nondegenerate quadratic form (matrix determinant $-c^2/4 \neq 0$), so the point is an ordinary node, multiplicity 2 as a point of the surface, resolved by a single blow-up. Substituting $u = Tu', y = Ty'$: the y^2 term becomes $T^2y'^2$; the $-cTu$ term becomes $-cT^2u'$; every other term, by hypothesis, carries an explicit factor T^2 already, or has total degree ≥ 4 in (u, T, y) and hence acquires an extra factor of T from each of the ≥ 2 occurrences of u, y being replaced by Tu', Ty' , giving total T -order ≥ 2 in every case, matching or exceeding the leading T^2 . Dividing the total transform by T^2 (the point’s multiplicity) and restricting to $T = 0$: the y'^2 and $-cu'$ terms survive as $y'^2 - cu'$, while every other term (carrying T -order strictly greater than 2 before division, by the degree count above) vanishes at $T = 0$. The restriction $y'^2 - cu' = 0$ is linear in u' , hence reduced (not a perfect square), so T vanishes to order exactly 1 transverse to the exceptional curve: fiber-multiplicity 1. $\square$

Verification at all nine points. In each case below, the local equation is obtained by the same two-step blow-up procedure as in Lemma 5.3’s proof sketch (first blow-up $X = Tx_1, Y = Ty_1$ or its $T = -1, \infty$ analogues; second blow-up, where needed, centered at the remaining singular point of the first exceptional curve), and the resulting tangent cone is displayed explicitly, verifying it has the form required by Lemma C.1 (constant $c = \pm 1 \neq 0$ throughout – never the degenerate case $c = 0$, which would signal a further, unresolved spine node).

Fiber, point	Local coordinate	Tangent cone (form $y^2 - cTu$)
$T = 0, P_1$	$x_1 = 0$ (level 1)	$y_1^2 - Tx_1$
$T = 0, P_2$	$x_3 = -1$ (level 2)	$y_3^2 + Tx_4$
$T = 0, P_3$	$x_3 = -2$ (level 2)	$y_3^2 - Tx_4$
$T = -1, P_2$	$x_1 = 1$ (level 1)	$y_1^2 - sx_2$
$T = -1, P_1$	$x_3 = 0$ (level 2)	$y_3^2 - sx_4$
$T = -1, P_3$	$x_3 = 1$ (level 2)	$y_3^2 + sx_4$
$T = \infty, P_3$	$x_1 = -1$ (level 1)	$y_1^2 - Sx_2$
$T = \infty, P_1$	$x_3 = 0$ (level 2)	$y_3^2 - Sx_4$
$T = \infty, P_2$	$x_3 = -1$ (level 2)	$y_3^2 + Sx_4$

($s = T + 1$ at $T = -1$; $S = 1/T$ at $T = \infty$; “level 1”/“level 2” indicate whether the point is resolved at the first blow-up or requires the second, as in Lemma 5.3’s proof sketch.) Every entry has $c = \pm 1 \neq 0$, so Lemma C.1 applies at each of the nine points, giving fiber-multiplicity exactly 1 throughout – the computation is genuine and separate at each point (the $T = -1$ and $T = \infty$ columns are *not* obtained by transporting the $T = 0$ column through any symmetry, per Remark 5.10; the numerical coincidences of sign and magnitude across rows reflect the shared *structural* pattern proved once in Lemma C.1, not a shared derivation).

Lemma C.2 (Complete multiplicity verification). *Every one of the twelve marked components (the four legs met by O, P_1, P_2, P_3 , at each of the three I_2^* fibers) has fiber-multiplicity exactly 1.*

Proof. For P_1, P_2, P_3 at each fiber: the nine tangent-cone computations above, via Lemma C.1. For O at each fiber: O is the point at infinity of the Weierstrass model, never touched by any of

the blow-ups resolving the finite singular points above (those blow-ups are centered at points with finite affine coordinates); by Lemma 5.5, *any* section – O included – meets exactly one component, necessarily of multiplicity 1, as a direct consequence of $\sigma \cdot F = 1$. This is a different, and in fact more elementary, argument than the tangent-cone computation used for P_1, P_2, P_3 : it requires no local blow-up data at all, only the global intersection-number fact already established in Lemma 5.5. $\square$

Independent chart cross-check. As a further check (not required by Lemma C.1’s proof, which is already complete and self-contained, but included as an additional verification), the $T = 0$, P_1 -leg computation was redone in the *other* available blow-up chart (X_4 -dominant: $T = X_4 t_2$, $Y_4 = X_4 y_2$, instead of T -dominant): this gives, at $X_4 = 0$, the restriction $y_2^2 - t_2 = 0$ – again reduced, confirming multiplicity 1 by a second, independent computational route at this one point.

APPENDIX D. GRAPH RIGIDITY AND THE NÉRON–SEVERI LATTICE, IN FULL

The graph. Vertices $L_1, L_2, S_1, S_2, S_3, L_3, L_4$; edges L_1 – S_1 , L_2 – S_1 , S_1 – S_2 , S_2 – S_3 , S_3 – L_3 , S_3 – L_4 (six edges, seven vertices – a tree). Multiplicities: 1 on each L_i , 2 on each S_j .

Rigidity, restated with every step labeled. *Step 1*: in a tree, a leaf’s neighbor is unique (a leaf has exactly one incident edge, by definition). L_1 is a leaf; its unique neighbor is S_1 . *Step 2*: if φ fixes L_1 (i.e. $\varphi(L_1) = L_1$), then φ maps L_1 ’s neighbor set to $\varphi(L_1)$ ’s neighbor set: $\varphi(\{S_1\}) = \{S_1\}$, i.e. $\varphi(S_1) = S_1$. *Step 3*: identically, fixing L_3 gives $\varphi(S_3) = S_3$. *Step 4*: S_2 is adjacent to exactly S_1 and S_3 (its only two tree-edges); no leg is adjacent to both spine ends (each leg has exactly one edge, to its own end’s spine node); hence S_2 is the *unique* vertex of the whole graph adjacent to both S_1 and S_3 simultaneously. *Step 5*: since $\varphi(S_1) = S_1$, $\varphi(S_3) = S_3$, and φ preserves adjacency, $\varphi(S_2)$ is also adjacent to both S_1 and S_3 ; by the uniqueness of Step 4, $\varphi(S_2) = S_2$. *Conclusion*: φ fixes $L_1, L_2, S_1, S_2, S_3, L_3, L_4$ – all seven vertices – given only that it fixes L_1, L_2, L_3, L_4 .

The 20 generators. O (zero section), F (general fiber class), and for each of the three I_2^* fibers, its six non-identity components (i.e. all seven components of that fiber except the one containing O). Count: $2 + 3 \times 6 = 2 + 18 = 20$.

Frobenius trace, derived. By Section 5 (summarized: all four legs of each fiber are pinned down as individually $\mathbb{Q}$ -rational by the torsion-section marking of Appendix C, and graph rigidity above then forces the spine components too), every one of the 20 generators is individually fixed by Frobenius at every good odd prime. Each such fixed algebraic (Tate) class contributes eigenvalue exactly p ; summing over 20 generators, $\text{Tr}(F_p \mid \text{NS}(X_{\text{III}})) = 20p$.

Discriminant, line by line. D_6 ’s root lattice discriminant: the Cartan matrix of D_6 is

$$C = \begin{pmatrix} 2 & -1 & 0 & 0 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 2 & -1 & -1 \\ 0 & 0 & 0 & -1 & 2 & 0 \\ 0 & 0 & 0 & -1 & 0 & 2 \end{pmatrix},$$

(a chain of four nodes with a fork of two at the end, the standard D_6 Dynkin diagram), and $\det C = 4$ – computed directly by cofactor expansion (independently verified by machine computation of this 6×6 determinant, not merely quoted from a table). U ’s Gram matrix is $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$, determinant -1 . Hence

$$\text{disc}(\text{Triv}(X_{\text{III}})) = \text{disc}(U) \cdot \text{disc}(D_6)^3 = (-1) \cdot 4^3,$$

$$4^3 = 64, \quad (-1) \cdot 64 = -64, \quad |\text{disc}(\text{Triv}(X_{\text{III}}))| = 64.$$

The elliptic-surface discriminant formula relates the trivial lattice's saturation-index to the torsion subgroup:

$$|\text{disc NS}(X_{\text{III}})| = \frac{|\text{disc}(\text{Triv}(X_{\text{III}}))|}{|\text{MW}_{\text{tors}}|^2} = \frac{64}{4^2} = \frac{64}{16} = 4.$$

(Each factor here – $\text{disc}(U) = -1$, $\text{disc}(D_6) = 4$ shown via the explicit Cartan determinant above, $|\text{MW}_{\text{tors}}| = 4$ from Appendix C – is independently justified below rather than left as an unexplained “standard computation,” since the lattice-to-quadratic-form step immediately following this one has a non-obvious failure mode worth guarding against explicitly.)

APPENDIX E. TRANSCENDENTAL LATTICE, CM, AND MODULARITY, IN FULL

From $\rho = 20$, $|\text{disc NS}| = 4$ to the transcendental lattice. X_{III} is a K3 surface, so $b_2 = 22 = \rho(X_{\overline{\mathbb{Q}}}) + \text{rank } T(X_{\text{III}})$; with $\rho = 20$, $\text{rank } T(X_{\text{III}}) = 2$. Since NS and T are orthogonal complements inside the unimodular lattice $H^2(X, \mathbb{Z})$, $|\text{disc } T(X_{\text{III}})| = |\text{disc NS}(X_{\text{III}})| = 4$ (a standard fact about primitive sublattices of unimodular lattices, applied here, not re-derived). $T(X_{\text{III}})$ is even (as the transcendental lattice of any K3 surface is) and positive-definite (the signature of $T(X)$ for a singular K3 is $(2, 0)$). A rank-2 even positive-definite lattice with $|\det| = 4$ has Gram matrix $\begin{pmatrix} 2a & b \\ b & 2c \end{pmatrix}$ with $4ac - b^2 = 4$; the simplest (and, as shown next, forced) solution is $a = c = 1, b = 0$, giving

$$T(X_{\text{III}}) \cong \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix} = \text{diag}(2, 2).$$

Distinguishing the lattice from the quadratic form. The lattice $\text{diag}(2, 2)$ is *not* the same object as the quadratic form with coefficients $(2, 0, 2)$. The correct dictionary: for Gram matrix $\begin{pmatrix} 2a & b \\ b & 2c \end{pmatrix}$, the associated *primitive* binary quadratic form is $Q(x, y) = ax^2 + bxy + cy^2$ – using *half* the diagonal Gram entries, not the entries themselves. For $\text{diag}(2, 2)$: $2a = 2 \Rightarrow a = 1$; $2c = 2 \Rightarrow c = 1$; $b = 0$. Hence

$$(a, b, c) = (1, 0, 1), \quad b^2 - 4ac = 0 - 4(1)(1) = -4.$$

(The tempting but erroneous reading $(a, b, c) = (2, 0, 2)$ – taking the Gram matrix's diagonal entries directly, rather than halving them – would give $b^2 - 4ac = -16$, inconsistent with the independently established $|\text{disc } T| = 4$; this is exactly the kind of slip the halving convention above is stated explicitly to prevent.)

CM field. $D = -4$, so the associated field is $\mathbb{Q}(\sqrt{-4}) = \mathbb{Q}(2i) = \mathbb{Q}(i)$. The reduced-form condition $|b| \leq a \leq c$ holds for $(1, 0, 1)$ ($0 \leq 1 \leq 1$), and the class number $h(-4) = 1$ means $(1, 0, 1)$ is the *unique* reduced form of discriminant -4 – so this identification is forced, not merely consistent.

Cross-check against Classes I/II. [1] states $T(X) \cong \text{diag}(2, 4)$, discriminant 8. Applying the identical dictionary: $2a = 2 \Rightarrow a = 1$, $2c = 4 \Rightarrow c = 2$, $b = 0$, giving $(a, b, c) = (1, 0, 2)$, $D = 0 - 4(1)(2) = -8$, field $\mathbb{Q}(\sqrt{-8}) = \mathbb{Q}(\sqrt{-2})$ – matching their stated result exactly. This cross-check, performed using the *same* convention on a case with an independently known answer, is the strongest available confirmation that the dictionary used here (halved diagonal entries) is the correct one.

The newform. Livné's theorem (Theorem 6.1) applies (hypotheses verified in Section 6), guaranteeing a weight-3, CM-by- $\mathbb{Q}(i)$ newform of some 2-power level governing $T(X_{\text{III}})$. Search of the relevant space identifies the unique candidate with rational eigenvalues at level 16: $f = \eta^6(4z) = 16.3.c.a.$

Twist elimination, reconstructed.

- (1) *Candidate set*: bad reduction only at 2 (Section 6.4) forces any twist to be by a character unramified outside $\{2\}$; the fundamental discriminants supported only at 2 are exactly $1, -4, 8, -8$, giving $\{1, \chi(-1), \chi(2), \chi(-2)\}$ – four candidates, proved exhaustive *before* any numerical evaluation.
- (2) *Two effective candidates*: $a_p(f) = 0$ at $p \equiv 3 \pmod{4}$ (CM vanishing), so all four coincide there (uninformative); at $p \equiv 1 \pmod{4}$, $\chi(-1)(p) = 1$ collapses $\{1, \chi(-1)\}$ to $F_0(p) = a_p(f)$, and $\chi(-2)(p) = \chi(-1)(p)\chi(2)(p) = \chi(2)(p)$ collapses $\{\chi(2), \chi(-2)\}$ to $F_1(p) = \chi(2)(p)a_p(f)$.
- (3) $p = 5$ *discriminator*: $5 \equiv 1 \pmod{4}$; $\chi(2)(5) = -1$ (since $5 \not\equiv \pm 1 \pmod{8}$); $a_5(f) = -6$; so $F_0(5) = -6$, $F_1(5) = +6$ – genuinely different. Direct evaluation of the 25-term sum gives $T(5) = -6 = F_0(5) \neq F_1(5)$.
- (4) *Untwisted representation*: since $\{F_0, F_1\}$ was proved exhaustive (step 2) *before* step 3's computation, and step 3 refutes F_1 , only F_0 remains possible – a deductive elimination, not a pattern match. $\text{Tr}(F_p \mid T(X_{\text{III}})) = a_p(f)$ for every odd p .

Why one prime deductively suffices: this is a general feature of elimination arguments over a finite, exhausted hypothesis set – once step 1 proves the set is $\{1, \chi(-1), \chi(2), \chi(-2)\}$ and no more, and step 2 proves this collapses (on the only primes where anything is testable) to exactly two functions of p , a *single* prime at which those two functions disagree suffices to falsify one of them completely, by ordinary logic (not statistics) – no larger sample size adds rigor once the candidate set itself is proved closed.

APPENDIX F. POINT-COUNT LEDGER AND REPRODUCIBILITY

The Lefschetz decomposition, by cohomological degree. For the smooth projective K3 surface $X_{\text{III}}/\mathbb{F}_p$, the Grothendieck–Lefschetz trace formula gives

$$\#X_{\text{III}}(\mathbb{F}_p) = \sum_{i=0}^4 (-1)^i \text{Tr}(F_p \mid H_{\text{et}}^i(X_{\text{III}}, \mathbb{Q}_\ell)).$$

- H^0 : rank 1, Frobenius acts trivially, $\text{Tr} = 1$.
- H^1 : rank 0 for any K3 surface ($b_1 = 0$), contributes 0.
- H^2 : rank $b_2 = 22$, splitting as $\text{NS}(X_{\text{III}}) \otimes \mathbb{Q}_\ell \oplus T(X_{\text{III}}) \otimes \mathbb{Q}_\ell$ (ranks $20 + 2$), contributing $\text{Tr}(F_p \mid \text{NS}) + \text{Tr}(F_p \mid T) = 20p + a_p(f)$ (Appendices D, E).
- H^3 : rank 0 ($b_3 = 0$ for K3), contributes 0.
- H^4 : rank 1, Frobenius acts by p^2 (top-degree Tate twist), $\text{Tr} = p^2$.

Summing (all even-degree terms with $+$ sign, odd-degree with $-$ sign, but the odd terms vanish):

$$\#X_{\text{III}}(\mathbb{F}_p) = 1 + (20p + a_p(f)) + p^2 = 1 + p^2 + 20p + a_p(f).$$

The affine/projective correction, re-derived independently. Write $\#X_{\text{III}}(\mathbb{F}_p)$ instead as a direct sum over the fibration: $p - 2$ good finite fibers, each smooth projective elliptic curve contributing $p + 1 - a_T$ points; three I_2^* fibers, each (now that every component is proved individually $\mathbb{F}_p$ -rational, Appendix D) contributing $7(p + 1) - 6 = 7p + 1$ points. At the two finite bad affine fibers, $T(T + 1) = 0$ kills V_{III} 's right-hand side identically, giving exactly p affine points each (Lemma 8.1); for good T , the affine count is $p - a_T$ (Lemma 8.2). Combining $\#V_{\text{III}}(\mathbb{F}_p) = p^2 + T(p)$ (definition) with these two lemmas gives $\sum_{\text{good } T} a_T = -T(p)$, and substituting into the fiber-by-fiber sum for $\#X_{\text{III}}(\mathbb{F}_p)$ reproduces

$$\#X_{\text{III}}(\mathbb{F}_p) = (p - 2)(p + 1) - \sum a_T + 3(7p + 1) = 1 + p^2 + 20p + T(p)$$

(Proposition 8.3, full derivation in Section 8). Equating the two independent expressions for $\#X_{\text{III}}(\mathbb{F}_p)$ (cohomological, this appendix; combinatorial, Section 8) gives $T(p) = a_p(f)$ – the same conclusion, reached by two structurally different routes, a further consistency check.

Reproducibility table (illustrative only – not part of the proof, which uses only the single exact value $p = 5$ deductively, per Appendix E). $T(p)$ is computed directly from Definition 2.1/Lemma 2.2; $a_p(16.3.c.a)$ is read from the LMFDB record for the newform 16.3.c.a (<https://www.lmfdb.org/ModularForm/GL2/Q/holomorphic/16/3/c/a/>), consistent with the q -expansion already quoted in Section 6, whose first listed coefficient, $a_5 = -6$, is the one value used deductively in the proof itself):

p	$T(p)$	$a_p(16.3.c.a)$	match
5	−6	−6	✓
13	10	10	✓
17	−30	−30	✓
29	42	42	✓

Scripts. The character-sum and lattice identities above are independently verified by scripts in the companion research record’s ‘explorations/class_iii/scripts/’ directory: the projectivization and involution identities, torsion completeness, the graph automorphism enumeration, the fiber-multiplicity checks, the twist elimination (‘class3_twist_elimination.py’), and the point-count ledger (‘class3_ledger_check.py’). None of these scripts is cited as a step in any proof above; each is cited only as an independent numerical check on a claim already proved by the preceding symbolic or deductive argument.

REFERENCES

- [1] E. De Jesús, *Finite-Field Hypercuboid Character Sums and a Discriminant-8 Singular K3 Surface*, Zenodo (2026), <https://doi.org/10.5281/zenodo.22711323>. The companion paper to this sequel; we follow its notation for the hypercuboid construction (§2) and its projectivization technique (Lemma 2.2) exactly.
- [2] S. Ahlgren, K. Ono, D. Penniston, *Zeta functions of an infinite family of K3 surfaces*, Amer. J. Math. **124** (2002), no. 2, 353–368. The source of the modular form 16.3.c.a = $\eta^6(4z)$, identified as the newform governing their own $\lambda = 8$ surface (see the Introduction); **not** claimed as new in this paper.
- [3] B. van Geemen, J. Top, *An isogeny of K3 surfaces*, Bull. London Math. Soc. **38** (2006), no. 2, 209–223, <https://doi.org/10.1112/S0024609306018170>; preprint [arXiv:math/0309272](https://arxiv.org/abs/math/0309272). Studies [2]’s family $X_t : y^2 = xz(x + 1)(z + 1)(x + zt)$; records that the general member has Néron–Severi rank 19, and that the member at $t = -1$ – the surface shown in Section 9.1 to be isomorphic over $\mathbb{Q}$ to X_{III} – has transcendental lattice of rank 2 and discriminant 4, hence $\langle 2 \rangle \oplus \langle 2 \rangle$, citing [4], and is by [12] the desingularised quotient of $E_i \times E_i$ with $E_i = \mathbb{C}/\mathbb{Z}[i]$.
- [4] U. Persson, *Double sextics and singular K-3 surfaces*, in *Algebraic Geometry, Sitges 1983*, Lecture Notes in Math. **1124**, Springer, 1985, 262–328, <https://doi.org/10.1007/BFb0075003>. The source [3] cites (p. 298) for the transcendental lattice of the $t = -1$ member; the systematic theory of double covers of $\mathbb{P}^2$ branched along a union of lines.
- [5] R. Wazir, *Arithmetic on elliptic threefolds*, Compositio Math. **140** (2004), no. 3, 567–580, <https://doi.org/10.1112/S0010437X03000381>; preprint [arXiv:math/0112259](https://arxiv.org/abs/math/0112259). In the enumeration of fiber types: for type I_n^* the number of rational components is $n + 5$ or $n + 3$, determined by the splitting of a quadratic polynomial. Prior art for the dichotomy underlying Lemma 5.7 (Remark 5.9).
- [6] D. Ulmer, *Elliptic curves over function fields*, in *Arithmetic of L-functions*, IAS/Park City Math. Ser. **18**, Amer. Math. Soc., 2011, 211–280, <https://doi.org/10.1090/pcms/018/09>; preprint [arXiv:1101.1939](https://arxiv.org/abs/1101.1939). Lecture 3, §6 defines split and non-split reduction for I_n^* by whether all geometric components are defined over the residue field, the non-split case leaving exactly one conjugate pair; cited alongside [5] in Remark 5.9.
- [7] K. Kodaira, *On compact analytic surfaces II–III*, Ann. of Math. **77** (1963), 563–626, and **78** (1963), 1–40. Source of the classification of singular fibers of elliptic surfaces used in Lemma 5.2.
- [8] A. Néron, *Modèles minimaux des variétés abéliennes sur les corps locaux et globaux*, Publ. Math. IHÉS **21** (1964), 5–128. Source of the Néron model and component-group theory underlying Lemma 5.5 and Proposition 4.5.
- [9] J. Tate, *Algorithm for determining the type of a singular fiber in an elliptic pencil*, in *Modular Functions of One Variable IV*, Lecture Notes in Math. **476**, Springer, 1975, 33–52. Used in Proposition 4.1 for the Kodaira-type determination criterion; see also Silverman, *Advanced Topics in the Arithmetic of Elliptic Curves*, Springer GTM 151, Ch. IV, for a standard modern exposition.
- [10] T. Shioda, *On the Mordell–Weil lattices*, Comment. Math. Univ. St. Paul. **39** (1990), no. 2, 211–240. Source of the Shioda–Tate formula (Proposition 4.2) and the torsion-injectivity theorem (Proposition 4.5).
- [11] R. Miranda, U. Persson, *Torsion groups of elliptic surfaces*, Compositio Math. **72** (1989), no. 3, 249–267. General theory of torsion groups of elliptic surfaces, cited alongside [10] in Proposition 4.5.

- [12] T. Shioda, H. Inose, *On singular K3 surfaces*, in *Complex Analysis and Algebraic Geometry*, Iwanami Shoten, Tokyo, 1977, 119–136. Source of the correspondence between rank-2 even lattices and imaginary-quadratic CM data used in Section 6.
- [13] R. Livné, *Motivic orthogonal two-dimensional representations of $\mathrm{Gal}(\overline{\mathbb{Q}}/\mathbb{Q})$* , Israel J. Math. **92** (1995), 149–156. Source of the singular-K3 modularity theorem (Theorem 6.1), applied here at discriminant -4 as it is applied in [1] at discriminant -8 .
- [14] T. Huber, Z.-G. Liu, et al., *On the vanishing of the coefficients of CM eta quotients*, preprint. Independent literature confirmation (outside [2]) that $\eta^6(4z) = \mathbf{16.3.c.a}$ is a CM newform by $\mathbb{Q}(i)$, there associated with the quartic Fermat K3 surface – a *third* variety, distinct from both X_8 and X_{III} , sharing the same trace formula (Remark 6.2).